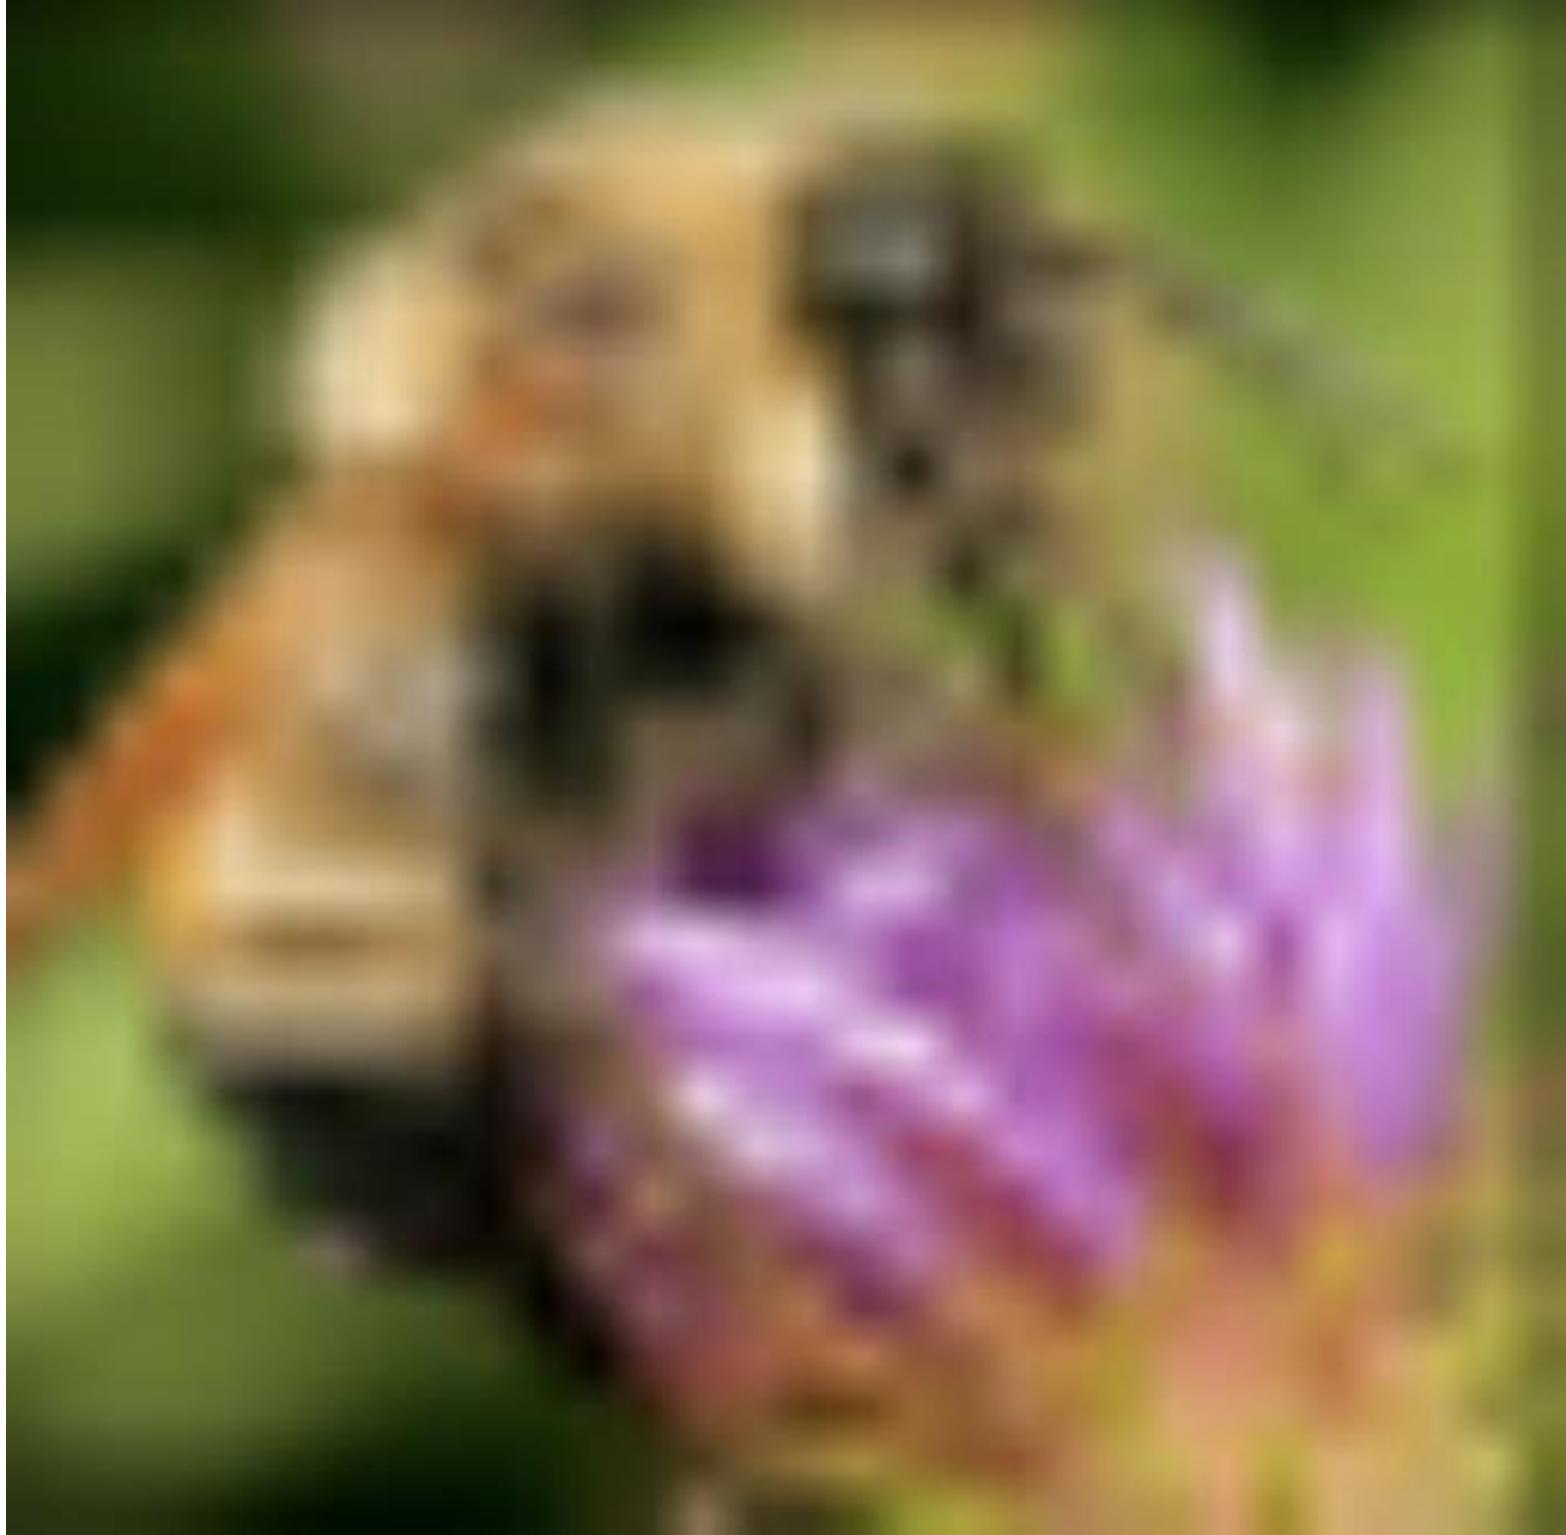


Q1.(a).

Code:

```
10     def OneDimLinearInterpolation(image, d, direction):
11
12         image_rows = image.shape[0]
13         image_cols = image.shape[1]
14
15         temp_image = np.zeros(shape=(image_rows, image_cols*d, 3))
16
17         rows = temp_image.shape[0]
18         cols = temp_image.shape[1]
19
20         result = []
21
22         if direction == 0:
23
24             for m in range(rows):
25                 temp = []
26                 for n in range(cols):
27                     i = n/d
28                     if i.is_integer():
29                         output = image[m, int(i)]
30                     else:
31                         x = i
32                         x_1 = math.floor(x)
33                         x_2 = math.ceil(x)
34                         if x_2 == image_cols:
35                             break
36                         else:
37                             r_1 = (x_2 - x) / (x_2 - x_1)
38                             r_2 = (x - x_1) / (x_2 - x_1)
39                             output = r_1 * image[m, x_1] + r_2 * image[m, x_2]
40                         temp.append(output)
41                 result.append(temp)
42
43         output_image = np.array(result)
44
45         if direction == 1:
46
47             rotated_image = np.rot90(image, -1)
48             temp_result = OneDimLinearInterpolation(rotated_image, d, 0)
49             output_image = np.rot90(temp_result, 1)
50
51     return output_image
52
```

1D Final Result: 685 × 677



First, I applied 1D interpolation, horizontally, then vertically, the result is above. If you do vertically first, then horizontally, it will also lead to the same result.





When  $d = 4$  (quadruple), the filter is:

```
[[ 0. , 0. , 0. , 0. , 0. , 0. , 0. , 0. , 0. ],
 [ 0. , 0.0625, 0.125 , 0.1875, 0.25 , 0.1875, 0.125 , 0.0625, 0. ],
 [ 0. , 0.125 , 0.25 , 0.375 , 0.5 , 0.375 , 0.25 , 0.125 , 0. ],
 [ 0. , 0.1875, 0.375 , 0.5625, 0.75 , 0.5625, 0.375 , 0.1875, 0. ],
 [ 0. , 0.25 , 0.5 , 0.75 , 1. , 0.75 , 0.5 , 0.25 , 0. ],
 [ 0. , 0.1875, 0.375 , 0.5625, 0.75 , 0.5625, 0.375 , 0.1875, 0. ],
 [ 0. , 0.125 , 0.25 , 0.375 , 0.5 , 0.375 , 0.25 , 0.125 , 0. ],
 [ 0. , 0.0625, 0.125 , 0.1875, 0.25 , 0.1875, 0.125 , 0.0625, 0. ],
 [ 0. , 0. , 0. , 0. , 0. , 0. , 0. , 0. , 0. ]]
```

### Q1.(b).

Suppose that we want to find the value of the unknown function  $f$  at the point  $(x, y)$ . It is assumed that we know the value of  $f$  at the four points  $Q_{11} = (x_1, y_1)$ ,  $Q_{12} = (x_1, y_2)$ ,  $Q_{21} = (x_2, y_1)$ , and  $Q_{22} = (x_2, y_2)$ .

We first do linear interpolation in the  $x$ -direction. This yields

$$f(x, y_1) \approx \frac{x_2 - x}{x_2 - x_1} f(Q_{11}) + \frac{x - x_1}{x_2 - x_1} f(Q_{21}),$$
$$f(x, y_2) \approx \frac{x_2 - x}{x_2 - x_1} f(Q_{12}) + \frac{x - x_1}{x_2 - x_1} f(Q_{22}).$$

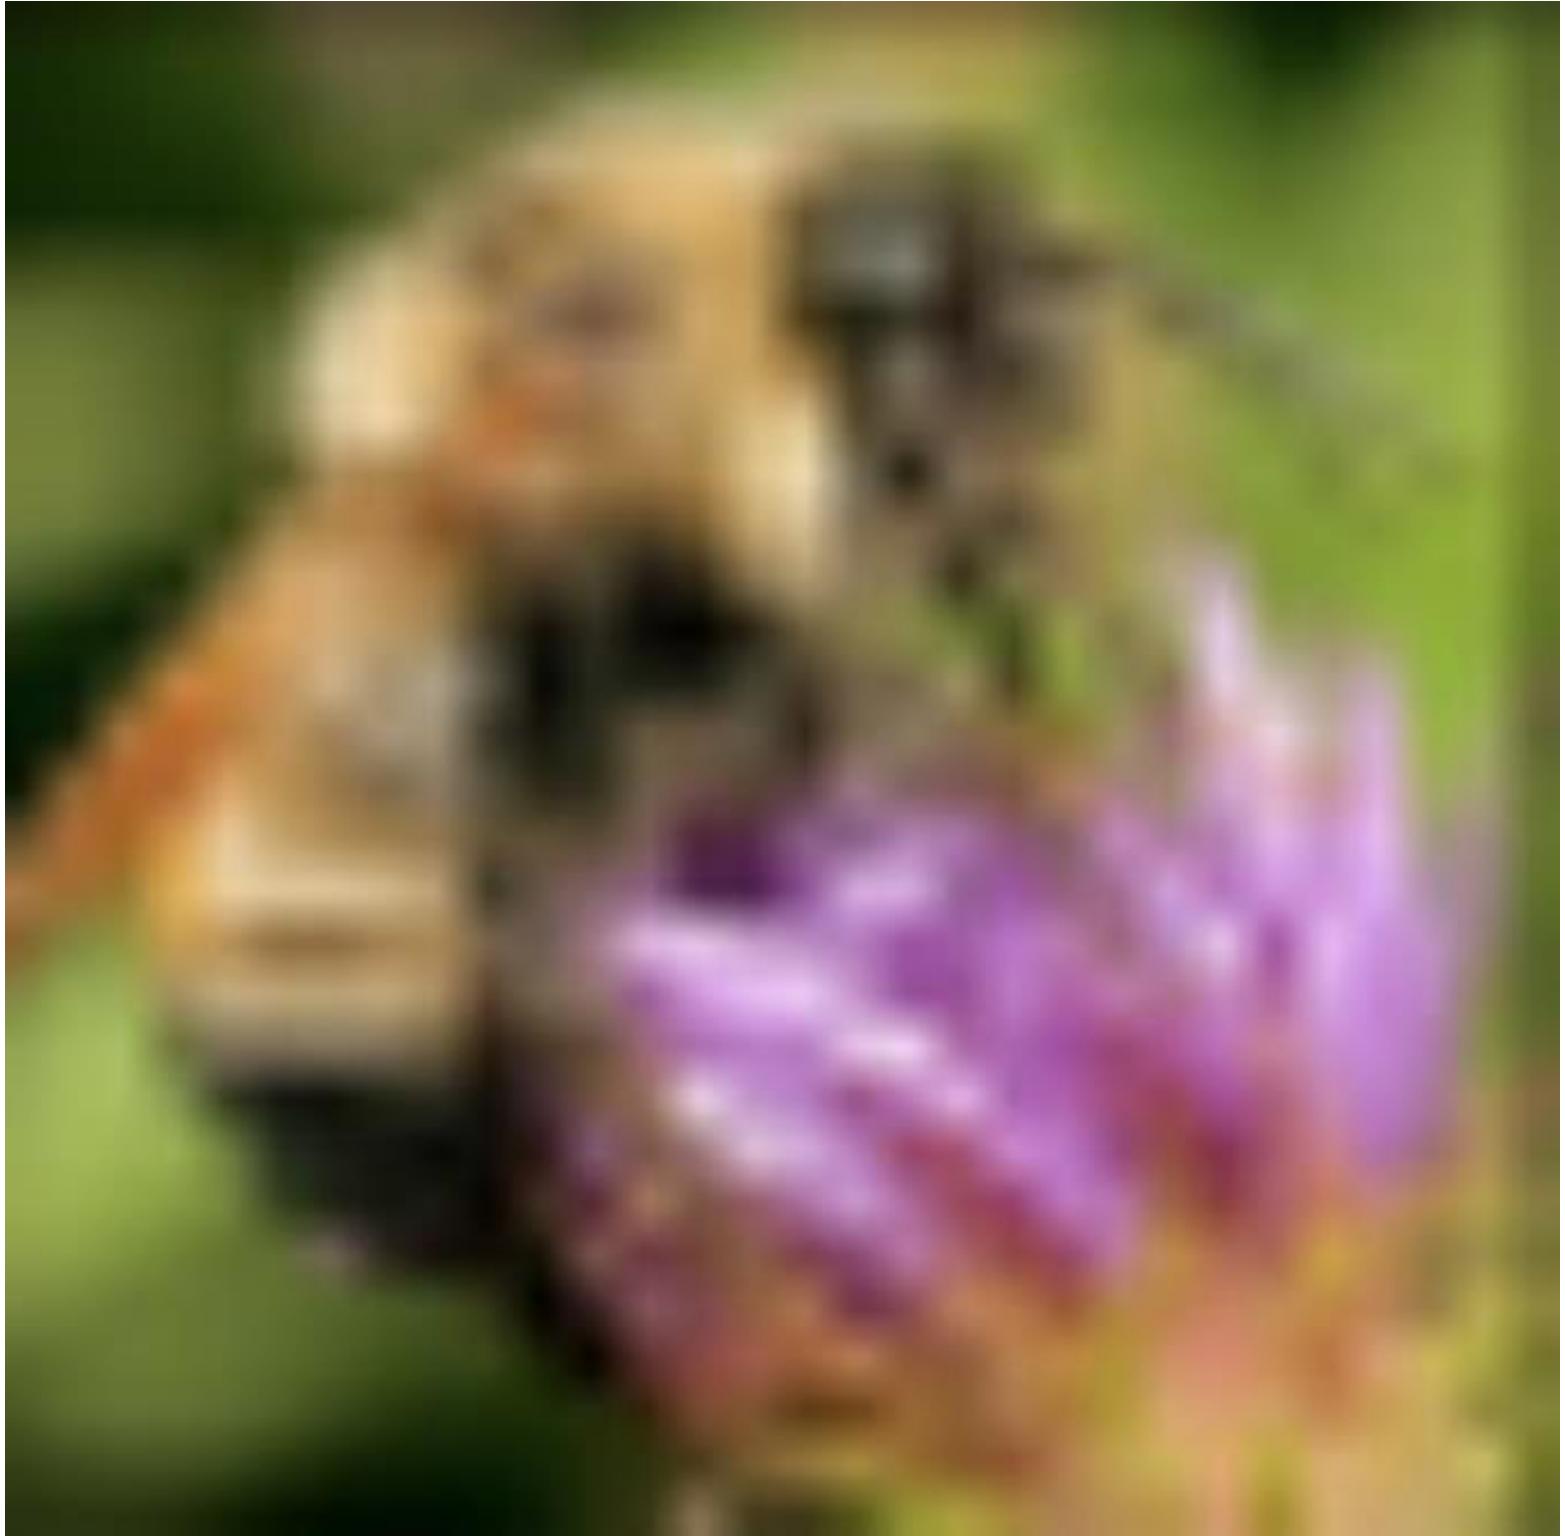
We proceed by interpolating in the  $y$ -direction to obtain the desired estimate:

$$\begin{aligned} f(x, y) &\approx \frac{y_2 - y}{y_2 - y_1} f(x, y_1) + \frac{y - y_1}{y_2 - y_1} f(x, y_2) \\ &= \frac{y_2 - y}{y_2 - y_1} \left( \frac{x_2 - x}{x_2 - x_1} f(Q_{11}) + \frac{x - x_1}{x_2 - x_1} f(Q_{21}) \right) + \frac{y - y_1}{y_2 - y_1} \left( \frac{x_2 - x}{x_2 - x_1} f(Q_{12}) + \frac{x - x_1}{x_2 - x_1} f(Q_{22}) \right) \\ &= \frac{1}{(x_2 - x_1)(y_2 - y_1)} (f(Q_{11})(x_2 - x)(y_2 - y) + f(Q_{21})(x - x_1)(y_2 - y) + f(Q_{12})(x_2 - x)(y - y_1) + f(Q_{22})(x - x_1)(y - y_1)) \\ &= \frac{1}{(x_2 - x_1)(y_2 - y_1)} [x_2 - x \quad x - x_1] \begin{bmatrix} f(Q_{11}) & f(Q_{12}) \\ f(Q_{21}) & f(Q_{22}) \end{bmatrix} \begin{bmatrix} y_2 - y \\ y - y_1 \end{bmatrix}. \end{aligned}$$

Note that we will arrive at the same result if the interpolation is done first along the  $y$  direction and then along the  $x$  direction.<sup>[1]</sup>

```
54 def TwoDimLinearInterpolation(image, d):
55
56     image_rows = image.shape[0]
57     image_cols = image.shape[1]
58
59     temp_image = np.zeros(shape=((image_rows - 1) * d + 1, (image_cols - 1) * d + 1, 3))
60
61     temp_image[:, :, ::d, ::d] = image.copy()
62
63     one_d_filter = []
64     for i in range(d):
65         one_d_filter.append(i/d)
66     for i in range(d, -1, -1):
67         one_d_filter.append(i/d)
68
69     one_d_filter = np.array(one_d_filter)
70
71     two_d_filter = np.matmul(np.transpose(np.array([one_d_filter])), np.array([one_d_filter]))
72
73     result = cv2.filter2D(temp_image, -1, two_d_filter)
74
75     return result
```

2D Final Result: 685 × 677



This result is the same as part (a).

Q2.(a).

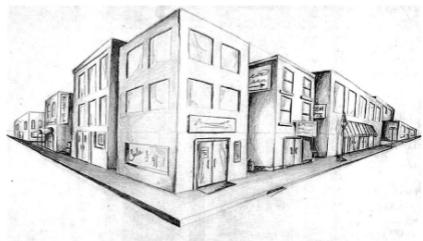
Test:

```
667  
668 # q2a_harris = harris(img_building, 0.05, 0.01)  
669 # q2a_harris2 = harris2(img_building, 0.05, 0.01)  
670  
671 # q2a_brown = brown(img_building)  
672 # q2a_brown2 = brown2(img_building)  
673  
674 # q2a_harris_suppression = harris(img_building, 0.05, 0.01, 1)  
675 # q2a_brown_suppression = brown(img_building, 5)
```

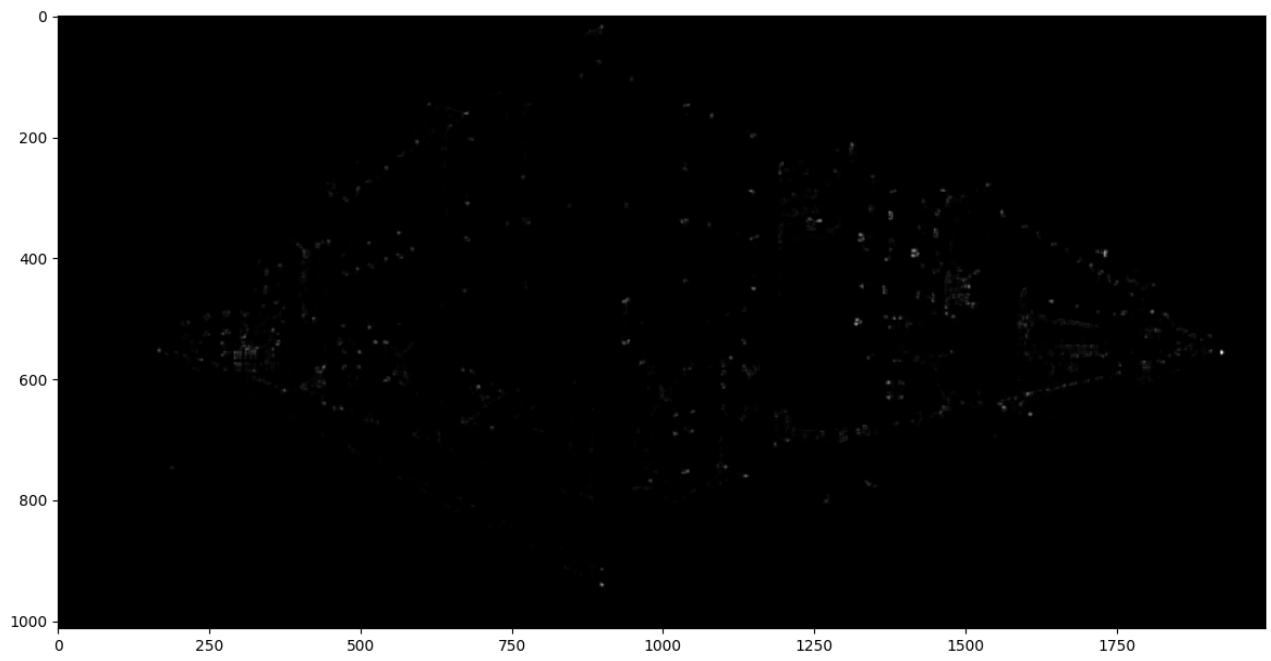
Harris:

```
78 def harris(image, alpha, threshold, radius):  
79  
80     gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)  
81     blur = cv2.GaussianBlur(gray, (5, 5), 7)  
82     ddepth = cv2.CV_64F  
83  
84     # STEP1: Compute gradient Ix and Iy  
85     Ix = cv2.Sobel(src=blur, ddepth=ddepth, dx=1, dy=0, ksize=5)  
86     Iy = cv2.Sobel(src=blur, ddepth=ddepth, dx=0, dy=1, ksize=5)  
87  
88     # STEP2: Compute Ix^2, Iy^2 and Ix*Iy  
89     IxIy = np.multiply(Ix, Iy)  
90     Ix2 = np.multiply(Ix, Ix)  
91     Iy2 = np.multiply(Iy, Iy)  
92  
93     # STEP3: Average(Gaussian)  
94     Ix2_blur = cv2.GaussianBlur(Ix2, (7, 7), 10)  
95     Iy2_blur = cv2.GaussianBlur(Iy2, (7, 7), 10)  
96     IxIy_blur = cv2.GaussianBlur(IxIy, (7, 7), 10)  
97  
98     # STEP4: Compute R  
99     det = np.multiply(Ix2_blur, Iy2_blur) - np.multiply(IxIy_blur, IxIy_blur)  
100    trace = Ix2_blur + Iy2_blur  
101    R = det - alpha * np.multiply(trace, trace)  
102  
103    # plt.subplot(1, 2, 1), plt.imshow(image), plt.axis('off')  
104    # plt.subplot(1, 2, 2), plt.imshow(R, cmap='gray'), plt.axis('off')  
105    # plt.imshow(R, cmap="gray")  
106    # plt.show()  
107  
108    #R = R * (R > (threshold * R.max())) * (R > 0)  
109    R[R < threshold] = 0  
110  
111    # plt.subplot(1, 2, 1), plt.imshow(image), plt.axis('off'), plt.show()  
112    # plt.subplot(1, 2, 2), plt.imshow(R, cmap='gray'), plt.axis('off'), plt.show()  
113    # plt.imshow(R, cmap="gray")  
114    # plt.show()  
115  
116    result = non_max_suppress(R, radius)  
117  
118    plt.imshow(result, cmap="gray")  
119    plt.show()  
120  
121    return 0
```

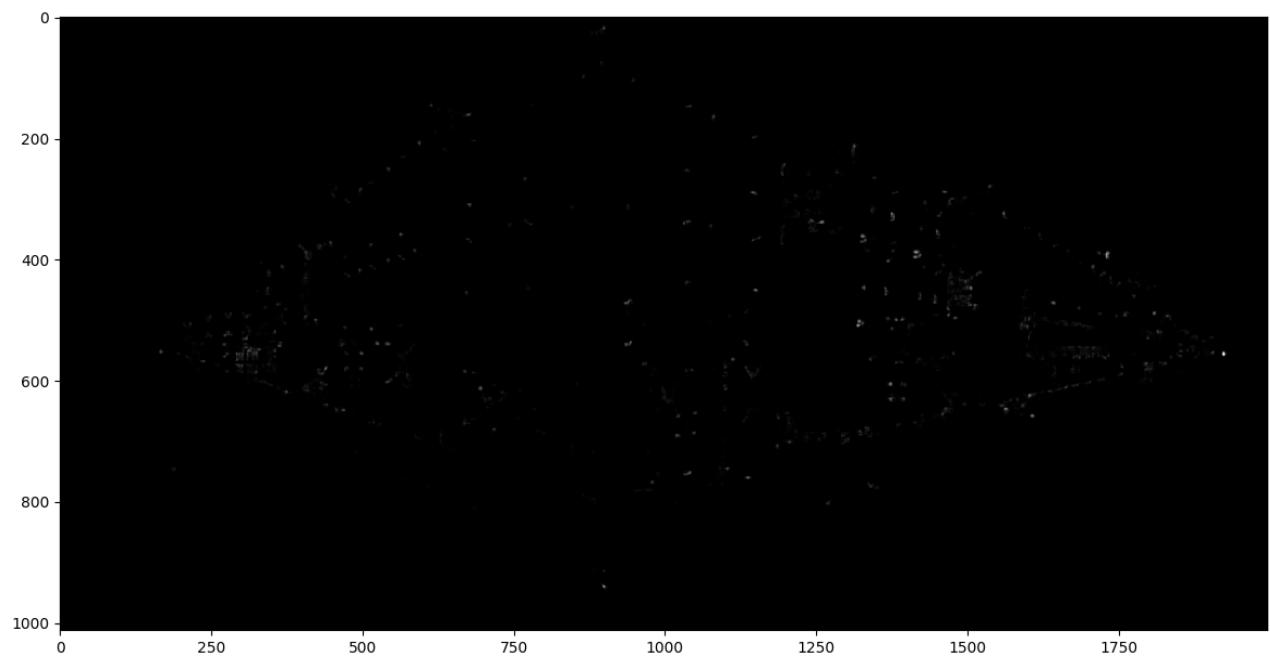
Harris Compare:



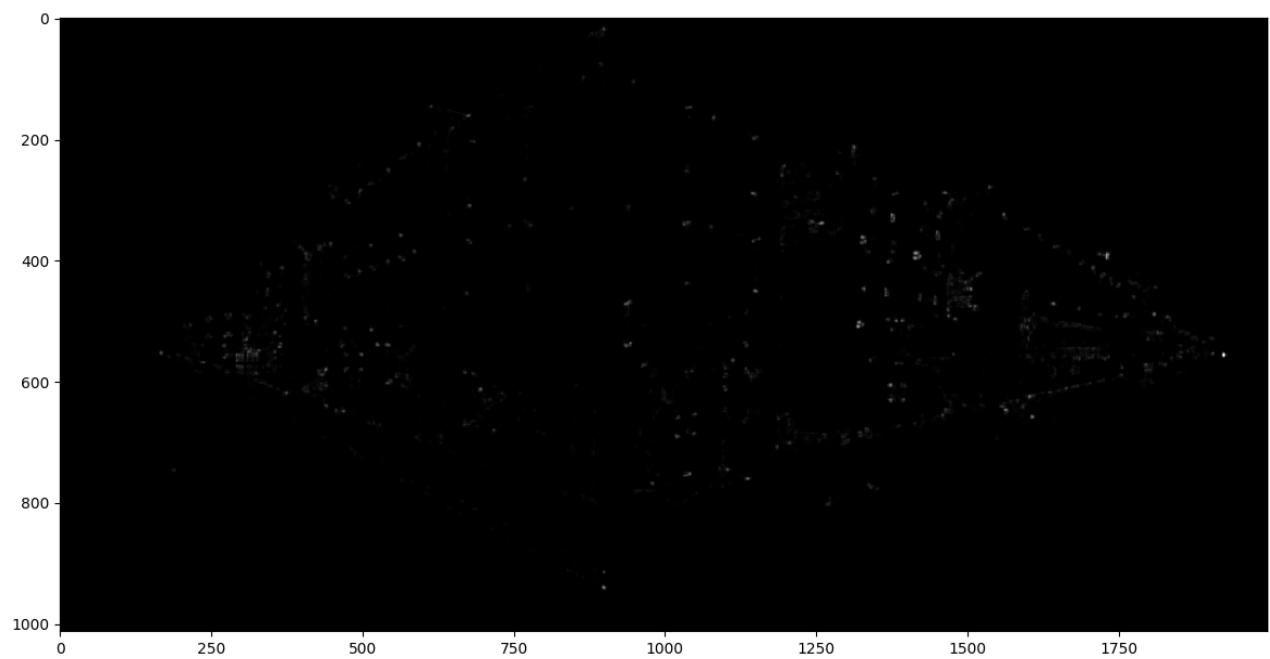
Harris Threshold:  
Alpha = 0.01



Alpha = 0.05

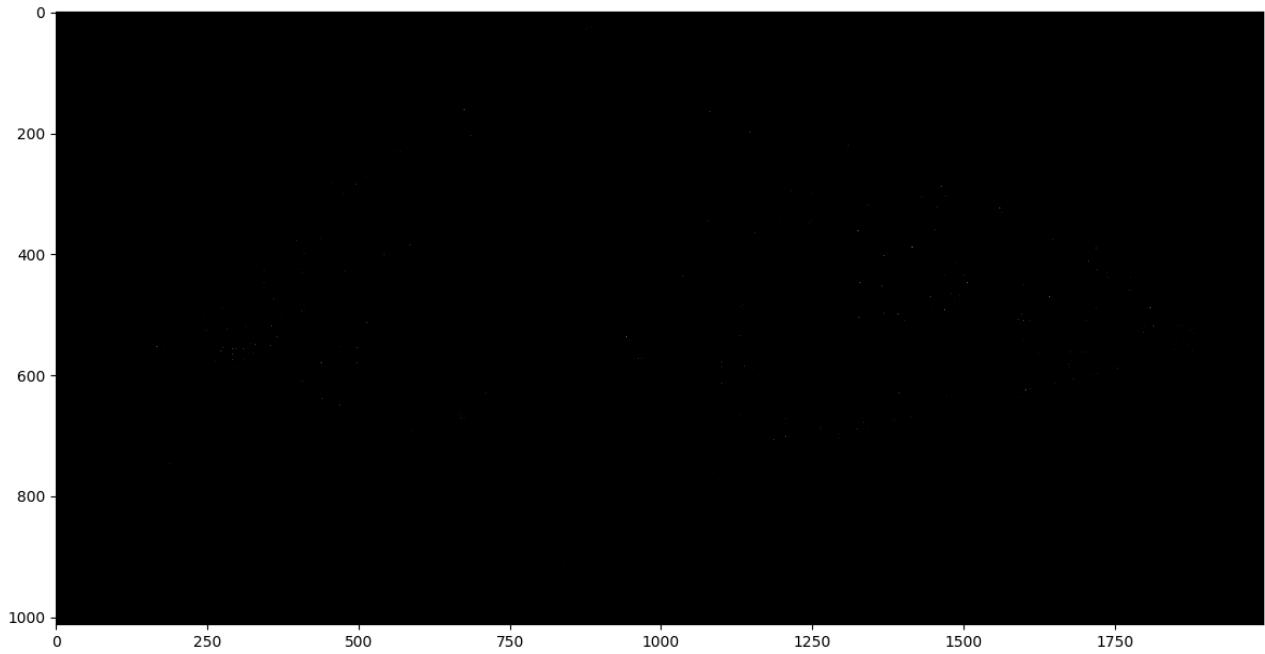


Alpha = 0.1

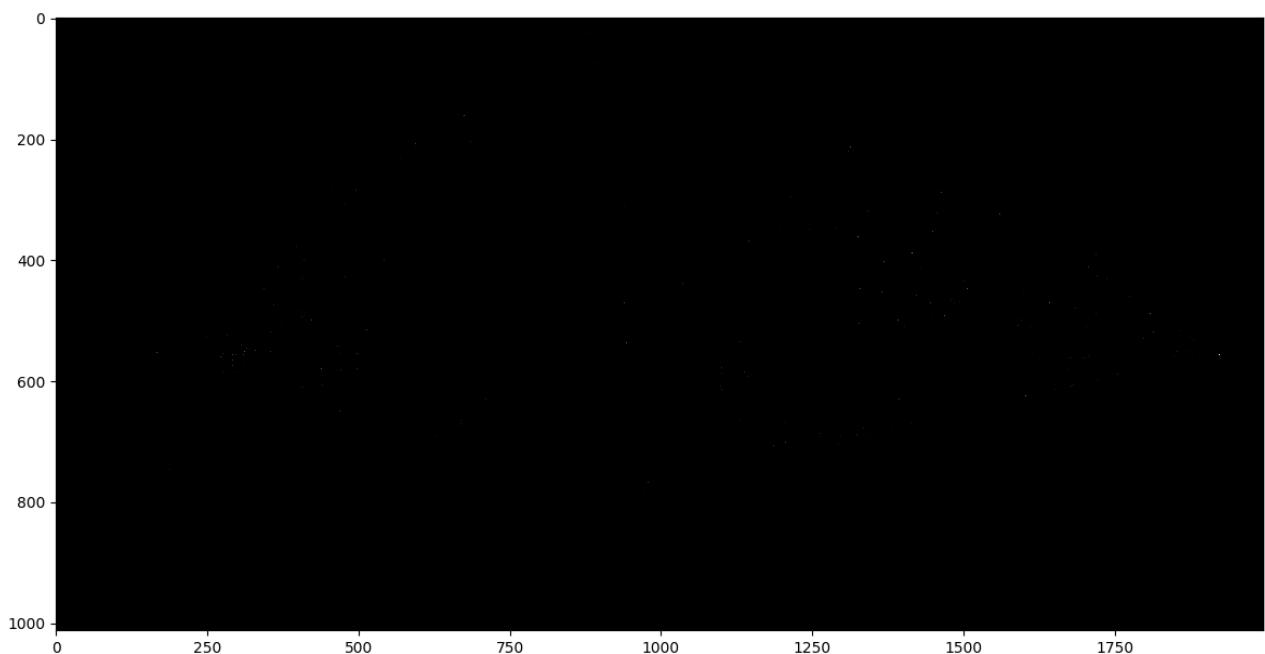


### Harris Non-maximal Suppression:

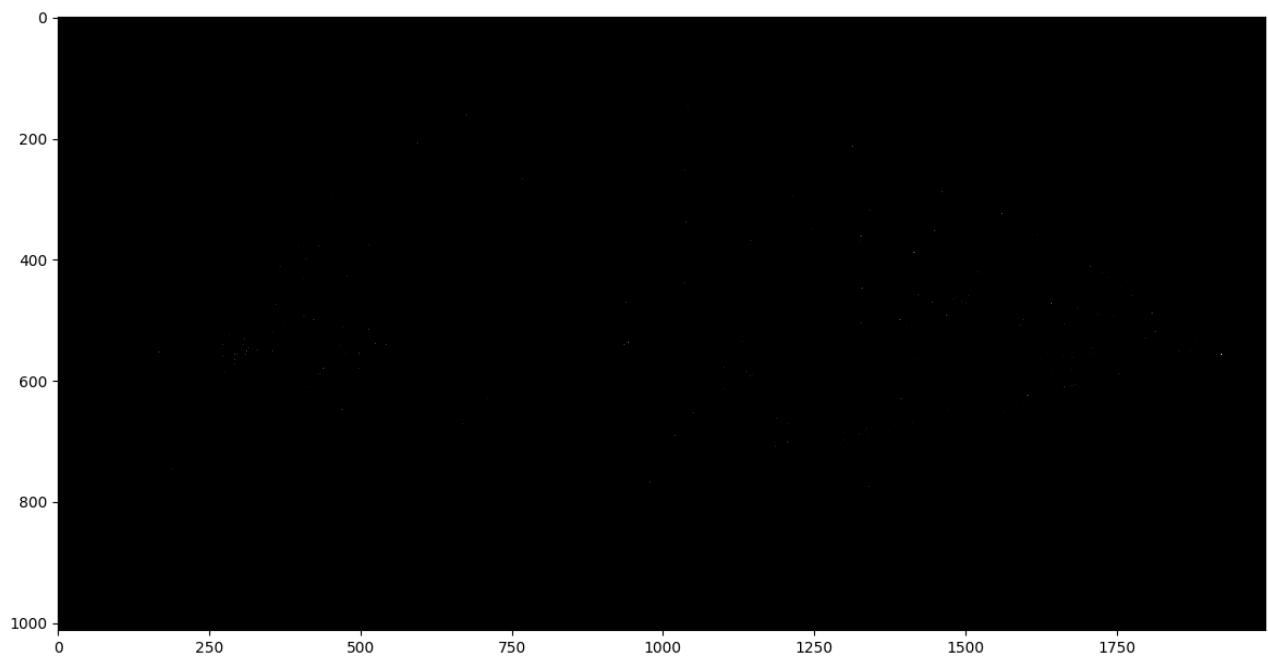
Alpha = 0.01



Alpha = 0.05



Alpha = 0.1

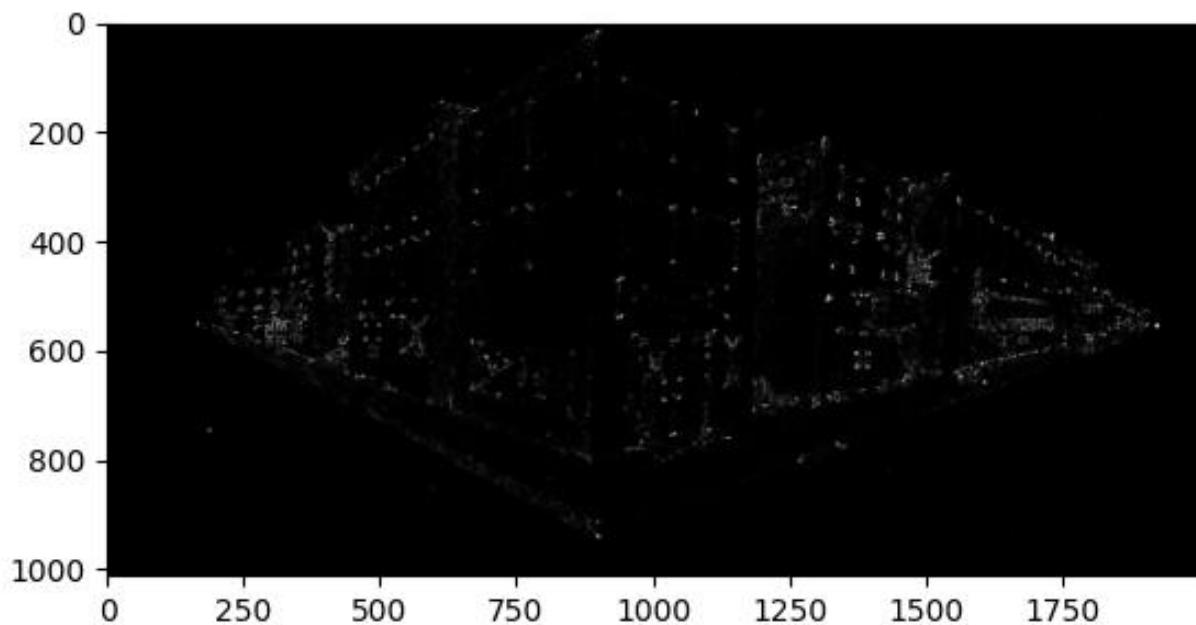


There is no big difference of these when alpha increasing little by little, the corners is slightly clearer.

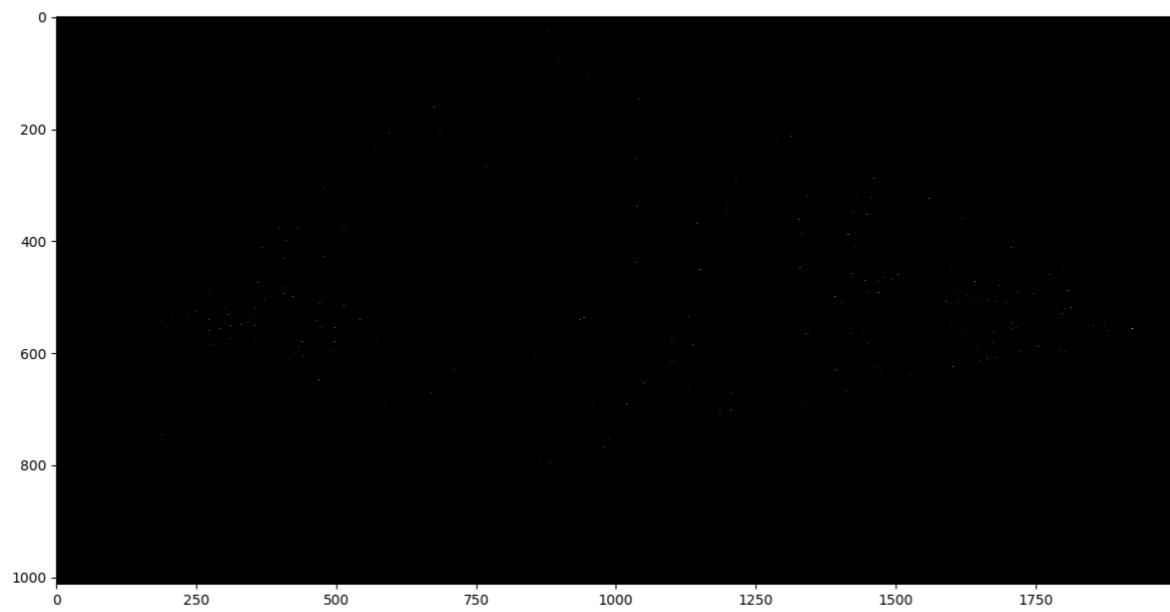
Brown:

```
104
105  def brown(image, radius):
106
107      gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
108      blur = cv2.GaussianBlur(gray, (5, 5), 7)
109      ddepth = cv2.CV_64F
110
111      # STEP1: Compute gradient Ix and Iy
112      Ix = cv2.Sobel(src=blur, ddepth=ddepth, dx=1, dy=0, ksize=5)
113      Iy = cv2.Sobel(src=blur, ddepth=ddepth, dx=0, dy=1, ksize=5)
114
115      # STEP2: Compute Ix^2, Iy^2 and Ix*Iy
116      IxIy = np.multiply(Ix, Iy)
117      Ix2 = np.multiply(Ix, Ix)
118      Iy2 = np.multiply(Iy, Iy)
119
120      # STEP3: Average(Gaussian)
121      Ix2_blur = cv2.GaussianBlur(Ix2, (7, 7), 10)
122      Iy2_blur = cv2.GaussianBlur(Iy2, (7, 7), 10)
123      IxIy_blur = cv2.GaussianBlur(IxIy, (7, 7), 10)
124
125      # STEP4: Compute R
126      det = np.multiply(Ix2_blur, Iy2_blur) - np.multiply(IxIy_blur, IxIy_blur)
127      trace = Ix2_blur + Iy2_blur
128      R = np.divide(det, trace+0.01)
129
130      # plt.subplot(1, 2, 1), plt.imshow(image), plt.axis('off'), plt.show()
131      # plt.subplot(1, 2, 2), plt.imshow(R, cmap='gray'), plt.axis('off'), plt.show()
132      # plt.imshow(R, cmap="gray")
133      # plt.show()
134
135      result = non_max_suppress(R, radius)
136
137      plt.imshow(result, cmap="gray")
138      plt.show()
139
140  return 0
```

Brown Threshold:



Brown non-maximal Suppression:



Harris without det and trace:

```
def harris2(image, alpha, threshold):

    gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
    blur = cv2.GaussianBlur(gray, (5, 5), 7)
    ddepth = cv2.CV_64F

    # STEP1: Compute gradient Ix and Iy
    Ix = cv2.Sobel(src=blur, ddepth=ddepth, dx=1, dy=0, ksize=5)
    Iy = cv2.Sobel(src=blur, ddepth=ddepth, dx=0, dy=1, ksize=5)

    # STEP2: Compute Ix^2, Iy^2 and Ix*Iy
    IxIy = np.multiply(Ix, Iy)
    Ix2 = np.multiply(Ix, Ix)
    Iy2 = np.multiply(Iy, Iy)

    # STEP3: Average(Gaussian)
    Ix2_blur = cv2.GaussianBlur(Ix2, (7, 7), 10)
    Iy2_blur = cv2.GaussianBlur(Iy2, (7, 7), 10)
    IxIy_blur = cv2.GaussianBlur(IxIy, (7, 7), 10)

    matrix = np.zeros((Ix2_blur.shape[0], Ix2_blur.shape[1], 2, 2))
    matrix[:, :, 0, 0] = Ix2_blur
    matrix[:, :, 0, 1] = IxIy_blur
    matrix[:, :, 1, 0] = IxIy_blur
    matrix[:, :, 1, 1] = Iy2_blur

    lambda_result = np.linalg.eigvals(matrix)
    lambda_0 = lambda_result[:, :, 0]
    lambda_1 = lambda_result[:, :, 1]

    # STEP4: Compute R
    R = lambda_0 * lambda_1 - alpha * (lambda_0 + lambda_1) ** 2
    R[R < threshold] = 0

    plt.subplot(1, 2, 1), plt.imshow(image), plt.axis('off'), plt.show()
    plt.subplot(1, 2, 2), plt.imshow(R, cmap='gray'), plt.axis('off'), plt.show()

    plt.imshow(R, cmap="gray")
    plt.show()

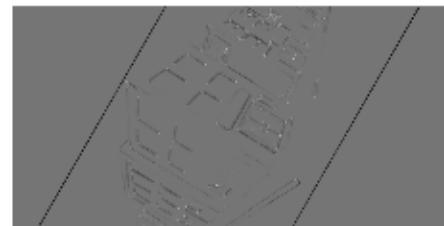
    return 0
```

Brown without det and trace:

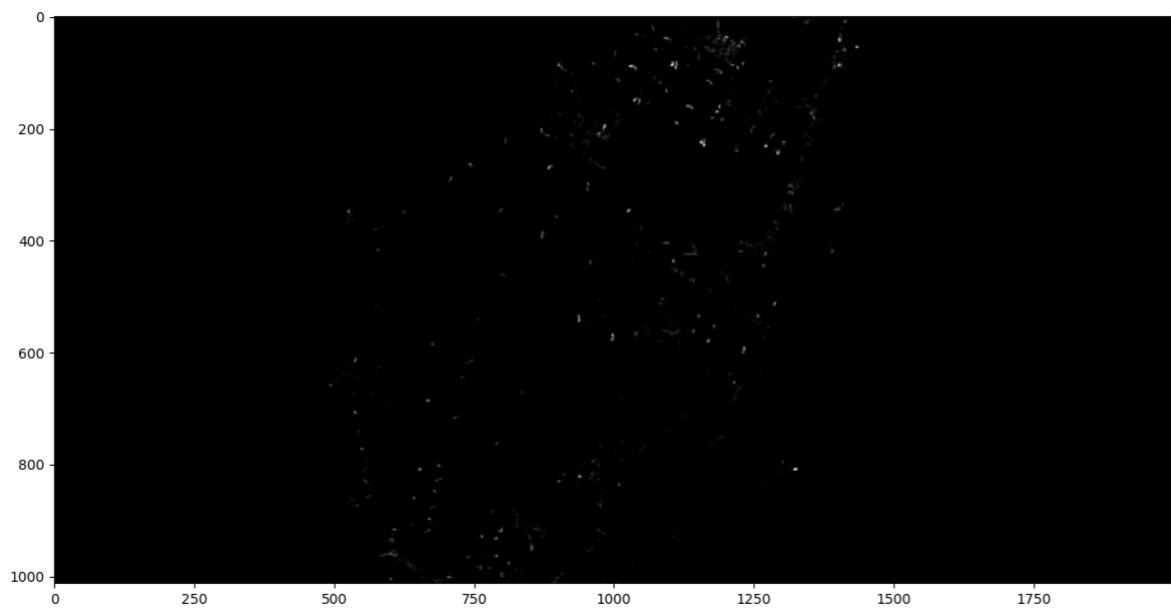
```
202
203     def brown2(image):
204
205         gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY)
206         blur = cv2.GaussianBlur(gray, (5, 5), 7)
207         ddepth = cv2.CV_64F
208
209         # STEP1: Compute gradient Ix and Iy
210         Ix = cv2.Sobel(src=blur, ddepth=ddepth, dx=1, dy=0, ksize=5)
211         Iy = cv2.Sobel(src=blur, ddepth=ddepth, dx=0, dy=1, ksize=5)
212
213         # STEP2: Compute Ix^2, Iy^2 and Ix*Iy
214         IxIy = np.multiply(Ix, Iy)
215         Ix2 = np.multiply(Ix, Ix)
216         Iy2 = np.multiply(Iy, Iy)
217
218         # STEP3: Average(Gaussian)
219         Ix2_blur = cv2.GaussianBlur(Ix2, (7, 7), 10)
220         Iy2_blur = cv2.GaussianBlur(Iy2, (7, 7), 10)
221         IxIy_blur = cv2.GaussianBlur(IxIy, (7, 7), 10)
222
223         matrix = np.zeros((Ix2_blur.shape[0], Ix2_blur.shape[1], 2, 2))
224         matrix[:, :, 0, 0] = Ix2_blur
225         matrix[:, :, 0, 1] = IxIy_blur
226         matrix[:, :, 1, 0] = IxIy_blur
227         matrix[:, :, 1, 1] = Iy2_blur
228
229         lambda_result = np.linalg.eigvals(matrix)
230         lambda_0 = lambda_result[:, :, 0]
231         lambda_1 = lambda_result[:, :, 1]
232
233         # STEP4: Compute R
234         R = np.divide(lambda_0 * lambda_1 , (lambda_0 + lambda_1)+0.01)
235
236         plt.subplot(1, 2, 1), plt.imshow(image), plt.axis('off'), plt.show()
237         plt.subplot(1, 2, 2), plt.imshow(R, cmap='gray'), plt.axis('off'), plt.show()
238
239         plt.imshow(R, cmap="gray")
240         plt.show()
241
242         return 0
243
```

Q2.(b).

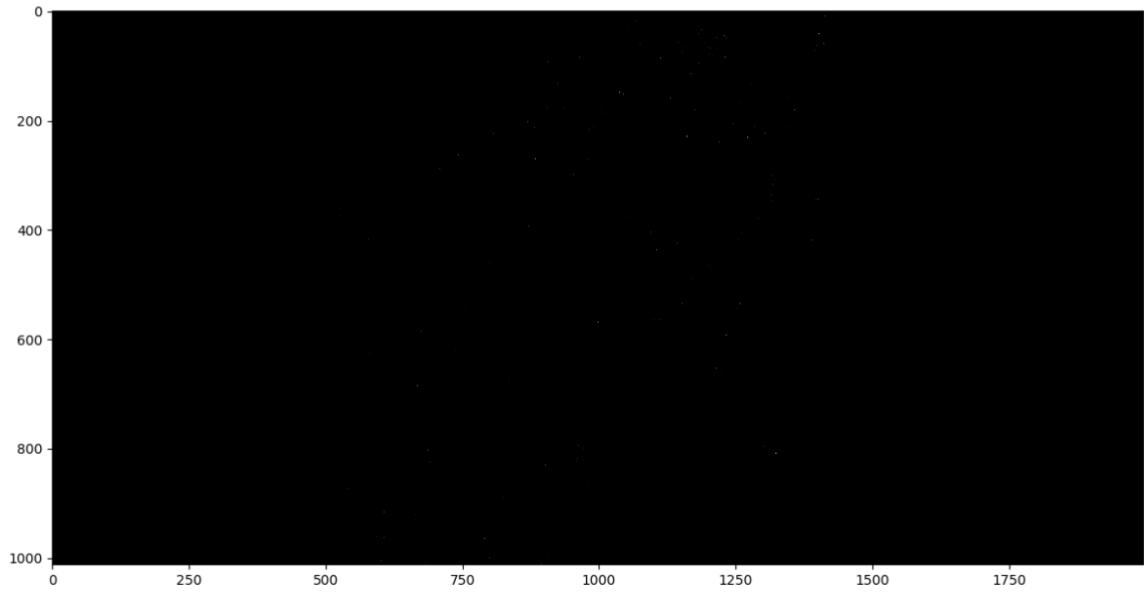
Rotated 60 degree:



Rotated Threshold:



## Rotated Non-maximal Suppression:

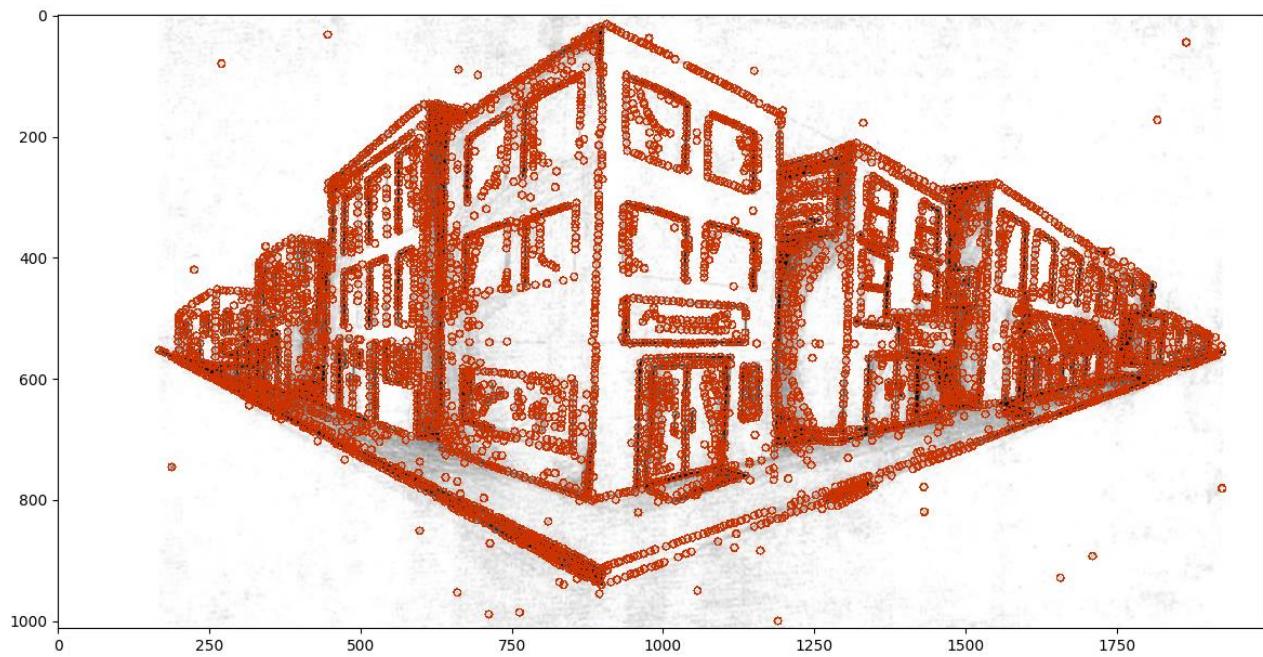


By observation, we can see the same corners can be still detected. Harris is rotation invariant because Second moment ellipse rotates but its shape (i.e. eigenvalues) remains the same.

## Q2.(c).

```
289     def LOG(image, levels, initial_sigma, scale, threshold):
290         gray = cv2.cvtColor(image, cv2.COLOR_BGR2GRAY).astype(np.float)
291
292         n = image.shape[0]
293         m = image.shape[1]
294
295         offset = 2
296         # construct pyramid
297         sigma = initial_sigma
298         sigma_list = [0]
299         increment = 1
300
301         pyramid = np.empty((levels + offset, n, m))
302
303         # pyramid = [np.zeros(gray.shape)]
304         for level in range(levels):
305             new_level = nd.gaussian_laplace(gray, sigma * increment)
306             pyramid[level + 1] = new_level
307             increment = increment * scale
308             sigma_list.append(sigma * increment)
309         # pyramid.append(np.zeros(gray.shape))
310
311         # Non Maximum Suppression
312         max_list = []
313         for level in range(1, levels):
314             # Construct upper current and lower matrix
315             # Current Level
316             cur_level = pyramid[level]
317             cur_level_mtx = np.zeros((n + offset, m + offset))
318             cur_level_mtx[1: n + 1, 1: m + 1] = cur_level[0: n, 0: m]
319             # Upper Level
320             upper_level = pyramid[level - 1]
321             up_level_mtx = np.zeros((n + offset, m + offset))
322             up_level_mtx[1: n + 1, 1: m + 1] = upper_level[0: n, 0: m]
323             # Lower Level
324             lower_level = pyramid[level + 1]
325             low_level_mtx = np.zeros((n + offset, m + offset))
326             low_level_mtx[1: n + 1, 1: m + 1] = lower_level[0: n, 0: m]
327
328             for row in range(n):
329                 for col in range(m):
330                     # Prepare for filtering
331                     patch_size = 3
332                     c_p = (patch_size - 1) // 2
333                     # Get currently filtering matrices from upper, current, lower levels
334                     upper_mtx = np.array(up_level_mtx[row: row + patch_size, col: col + patch_size])
335                     cur_mtx = np.array(cur_level_mtx[row: row + patch_size, col: col + patch_size])
336                     lower_mtx = np.array(low_level_mtx[row: row + patch_size, col: col + patch_size])
337                     # Get max value form upper, current, lower levels
338                     up_cur_low_max = []
339                     up_cur_low_max.append(np.max(upper_mtx))
340                     up_cur_low_max.append(np.max(cur_mtx))
341                     up_cur_low_max.append(np.max(lower_mtx))
342                     max_value = np.max(up_cur_low_max)
343                     # Get current filtering center point value
344                     c_p_value = cur_mtx[c_p, c_p]
345                     # Threshold
346                     if max_value == c_p_value and threshold < c_p_value:
347                         max_list.append(((row, col), sigma_list[level]))
348
349             for max in max_list:
350                 cv2.circle(image, (max[0][1],
351                                   max[0][0]),
352                            int(np.round(max[1])) * 3,
353                            color=(123, 255, 222),
354                            thickness=2)
355
356             plt.imshow(image)
357             plt.show()
```

```
# LOG(img_building, levels=5, initial_sigma=1.6, scale=np.sqrt(2), threshold=5)
```



Q2.(d).

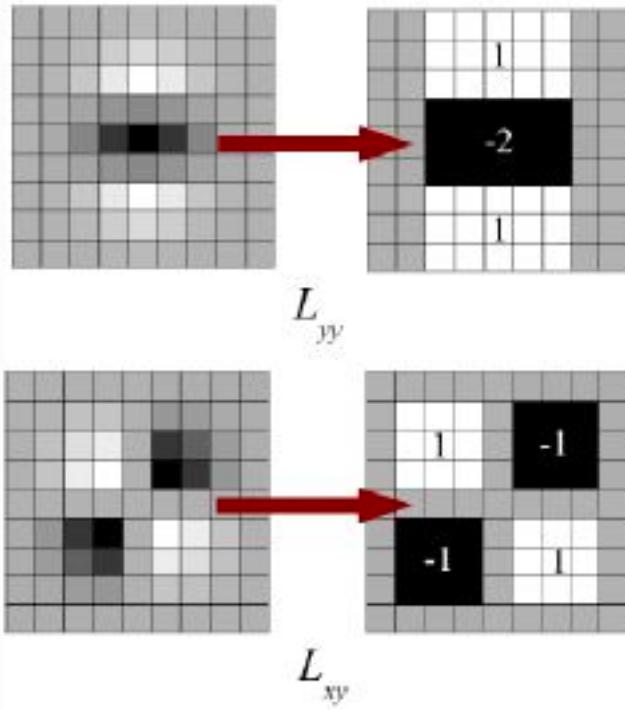
**SURF detector:** SURF stands for Speeded Up Robust Features which is inspired by SIFT and approaches corner detection via Laplacian of Gaussian. SURF uses an integer approximation of the determinant of Hessian blob detector using the precomputed integral image:

$$S(x, y) = \sum_{i=0}^x \sum_{j=0}^y I(i, j)$$

$$\text{Hessian matrix: } H(p, \sigma) = \begin{pmatrix} L_{xx}(p, \sigma) & L_{xy}(p, \sigma) \\ L_{yx}(p, \sigma) & L_{yy}(p, \sigma) \end{pmatrix}$$

The determinant of the Hessian matrix is evaluated as the measurement of the local changes around the point. For non-maximum suppression, choose the point with maximum determinant of Hessian in its neighbor as the key-point.

In SIFT, Lowe approximated Laplacian of Gaussian with Difference of Gaussian for finding scale-space. SURF goes a little further and approximates LoG with Box Filter. Below image shows a demonstration of such an approximation. One big advantage of this approximation is that, convolution with box filter can be easily calculated with the help of integral images. And it can be done in parallel for different scales. Also the SURF rely on determinant of Hessian matrix for both scale and location.



The creation of SURF descriptor takes place in two steps. The first step consists of fixing a reproducible orientation based on information from a circular region around the keypoint. Then, we construct a square region aligned to the selected orientation and extract the SURF descriptor from it.

1. Surf first calculate the Haar-wavelet responses in x and y-direction, and this in a circular neighborhood of radius  $6s$  around the keypoint, with  $s$  the scale at which the keypoint was detected. Also, the sampling step is scale dependent and chosen to be  $s$ , and the wavelet responses are computed at that current scale  $s$ . Accordingly, at high scales the size of the wavelets is big. Therefore integral images are used again for fast filtering.
2. Then we calculate the sum of vertical and horizontal wavelet responses in a scanning area, then change the scanning orientation ( $\text{add } \pi/3$ ), and re-calculate, until we find the orientation with largest sum value, this orientation is the main orientation of feature descriptor.
3. The first step consists of constructing a square region centered around the keypoint and oriented along the orientation we already got above. The size of this window is  $20s$ .
4. Then the region is split up regularly into smaller  $4 \times 4$  square sub-regions. For each sub-region, we compute a few simple features at  $5 \times 5$  regularly spaced sample points. For reasons of simplicity, we call  $\mathbf{dx}$  the Haar wavelet response in the horizontal direction and  $\mathbf{dy}$  the Haar wavelet response in the vertical direction (filter size  $2s$ ). To increase the robustness towards geometric deformations and localization errors, the responses  $\mathbf{dx}$  and  $\mathbf{dy}$  are first weighted with a Gaussian ( $\sigma = 3.3s$ ) centered at the keypoint.

Then, the wavelet responses  $\mathbf{dx}$  and  $\mathbf{dy}$  are summed up over each subregion and form a first set of entries to the feature vector. In order to bring in information about the polarity of the intensity changes, we also extract the sum of the absolute values of the responses,  $|\mathbf{dx}|$  and  $|\mathbf{dy}|$ . Hence, each sub-region has a four-dimensional descriptor vector  $\mathbf{v}$  for its underlying intensity structure  $\mathbf{V}$ . This results in a descriptor vector for all  $4 \times 4$  sub-regions of **length 64** (In **Sift**, our descriptor is the **128-D vector**, so this is part of the reason that SURF is faster than Sift).

Q3.(a).

Q3.

$$(a). G(x, y, \sigma) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

$$\frac{\partial G(x, y, \sigma)}{\partial x} = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}} \cdot \left(-\frac{2x}{2\sigma^2}\right)$$

$$= -\frac{x}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

$$\begin{aligned} -\frac{\partial^2 G}{\partial x^2} &= -\frac{1}{2\pi\sigma^4} \cdot \left(e^{-\frac{x^2+y^2}{2\sigma^2}} + e^{-\frac{x^2+y^2}{2\sigma^2}} \cdot -\frac{x^2}{\sigma^2}\right) \\ &= -\frac{1}{2\pi\sigma^4} \cdot e^{-\frac{x^2+y^2}{2\sigma^2}} \left(1 - \frac{x^2}{\sigma^2}\right) \end{aligned}$$

$$! \quad \frac{\partial^2 G}{\partial y^2} = -\frac{1}{2\pi\sigma^4} \cdot e^{-\frac{x^2+y^2}{2\sigma^2}} \left(1 - \frac{y^2}{\sigma^2}\right)$$

$$\therefore \nabla^2 G(x, y, \sigma) = \frac{\partial^2 G}{\partial x^2} + \frac{\partial^2 G}{\partial y^2} = -\frac{1}{2\pi\sigma^4} e^{-\frac{x^2+y^2}{2\sigma^2}} \left(2 - \frac{x^2+y^2}{\sigma^2}\right)$$

LoG is not separable.

LoG cannot be rewritten to a product of two linear filters because it has  $(2 - \frac{x^2+y^2}{\sigma^2})$  term in closed-form expression.

Also, like this example:  $\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix} \leftarrow \text{This LoG filter ranked as 2,}$

Q3.(b).

Q3(b). By SIFT paper from Lec 8, the scale-normalized  $\text{LoG } \sigma^2 \nabla^2 G$  has maxima & minima produce most stable features.

Heat diffusion equation:  $\frac{\partial G}{\partial \sigma} = \sigma \nabla^2 G$

So we can get:  $\sigma \nabla^2 G = \frac{\partial G}{\partial \sigma} \approx \frac{G(x, y, k\sigma) - G(x, y, \sigma)}{k\sigma - \sigma}$

( $\nabla^2 G$  can be computed from finite difference approximation to  $\frac{\partial G}{\partial \sigma}$ ) (Scale at  $k\sigma$  and  $\sigma$ )

Therefore,  $\text{D}_\sigma G$  (left hand side) can approximation to scale-normalized  $\text{LoG } \sigma^2 \nabla^2 G$ .  $G(x, y, k\sigma) - G(x, y, \sigma) \approx (k-1) \sigma^2 \nabla^2 G$

For  $\text{D}_\sigma G$ , let  $\sigma_1 = k\sigma$ ,  $\sigma_2 = \sigma$ .

We want  $\sigma_1 \approx \sigma_2 \Rightarrow k\sigma \approx \sigma \Rightarrow k \approx 1$

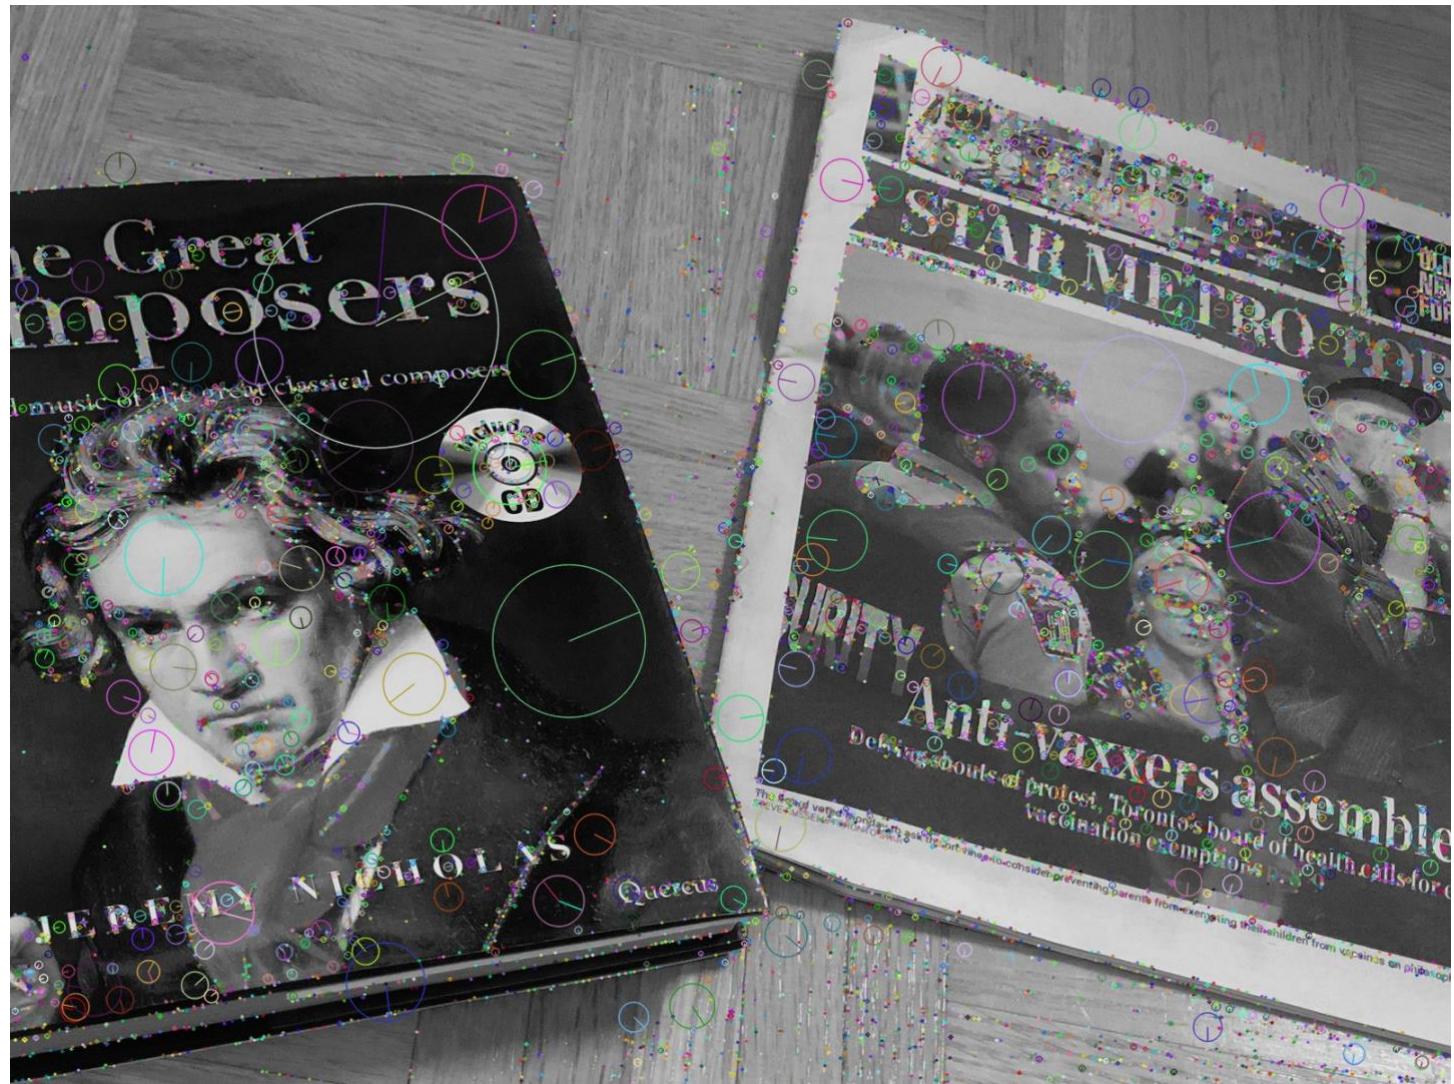
When  $k \approx 1$ , approximation error  $\Rightarrow 0$ .

This maximizes the approximation to  $\text{LoG } \sigma^2 \nabla^2 G$

Q4.(a).

```
360 def sift_extract(image):
361     img = cv2.imread(image)
362     gray_img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
363
364     # Initialized sift
365     sift = cv2.xfeatures2d.SIFT_create()
366     flag = cv2.DRAW_MATCHES_FLAGS_DRAW_RICH_KEYPOINTS
367
368     # Get keypoints and features
369     keypoints, features = sift.detectAndCompute(gray_img, None)
370     result = cv2.drawKeypoints(gray_img, keypoints, flags=flag, outImage=None)
371
372     cv2.imwrite('sift_keypoints.jpg', result)
373     return keypoints, features
```





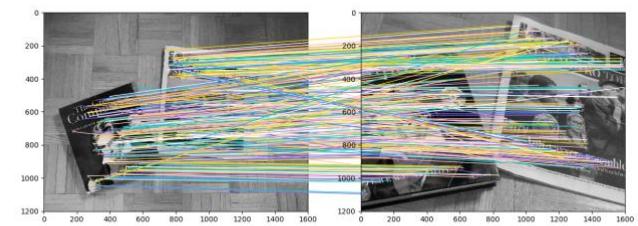
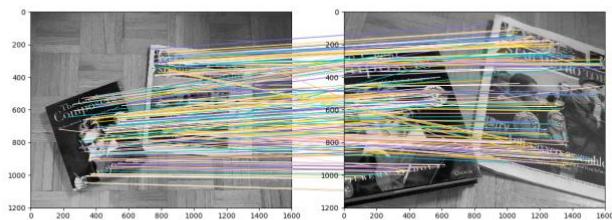
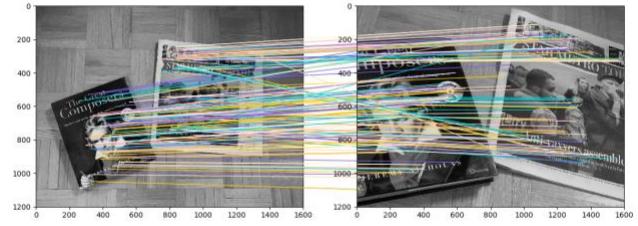
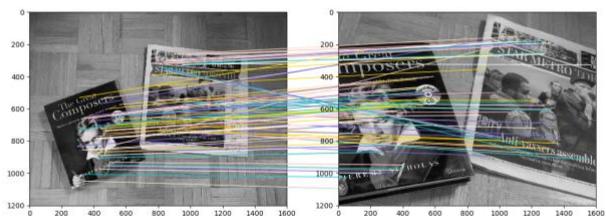
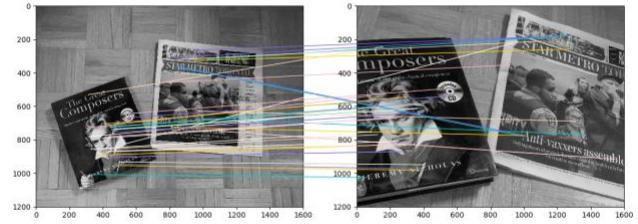
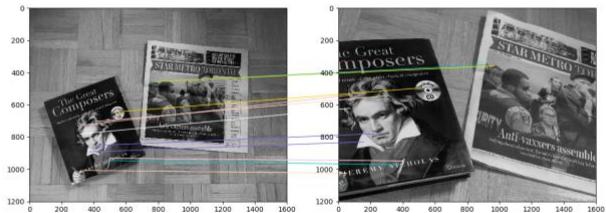
## Q4.(b).

```
376     def sift_matching(image1, image2, ratio_threshold, norm_level):
377         img1 = cv2.imread(image1)
378         img2 = cv2.imread(image2)
379
380         gray1 = cv2.cvtColor(img1, cv2.COLOR_BGR2GRAY)
381         gray2 = cv2.cvtColor(img2, cv2.COLOR_BGR2GRAY)
382
383         kp_1, des_1 = sift_extract(image1)
384         kp_2, des_2 = sift_extract(image2)
385
386         candidate_keypoint_location = []
387         candidate_distance = []
388
389         num_des_1 = np.shape(des_1)[0]
390         num_des_2 = np.shape(des_2)[0]
391
392         num_matching_keypoint = num_des_1
393
394         # Compare all keypoints
395         for i in range(num_des_1):
396             distance = []
397             for j in range(num_des_2):
398                 euclidean_dis = euclidean_distance(des_1[i], des_2[j], norm_level)
399                 distance.append(euclidean_dis)
400             # find two smallest distance match
401             two_smallest_distance = heapq.nsmallest(2, distance)
402             candidate_distance.append(two_smallest_distance)
403             # find minimum distance index at keypoint 2
404             np_distance = np.array(distance)
405             min_dis_index = np.unravel_index(np.argmin(np_distance, axis=None), np_distance.shape)[0]
406             candidate_keypoint_location.append(min_dis_index)
407
408
409             keypoint1 = []
410             keypoint2 = []
411             keypoints_pairs = {}
412             count = 0
413
414
415             for i in range(num_matching_keypoint):
416                 ratio = np.true_divide(candidate_distance[i][0], candidate_distance[i][1])
417                 if ratio < ratio_threshold:
418                     count = count + 1
419                     kp1 = kp_1[i]
420                     kp2 = kp_2[candidate_keypoint_location[i]]
421                     keypoint1.append(kp1)
422                     keypoint2.append(kp2)
423                     keypoints_pairs[(kp1, kp2)] = ratio
424
425             sorted_keypointes = sorted(keypoints_pairs.items(), key=lambda kv: kv[1])
426
427             print("number of matching = ", count)
428
429             top_ten = sorted_keypointes[0:9]
430             print(top_ten)
431
432             flag = cv2.DRAW_MATCHES_FLAGS_DRAW_RICH_KEYPOINTS
433
434             result_img1 = cv2.drawKeypoints(gray1, keypoint1, outImage=None, flags=flag)
435             result_img2 = cv2.drawKeypoints(gray2, keypoint2, outImage=None, flags=flag)
436             cv2.imwrite("sample1_keypoints_final.jpg", result_img1)
437             cv2.imwrite("sample2_keypoints_final.jpg", result_img2)
438             show_matching(gray1, keypoint1, gray2, keypoint2)
439
440             return top_ten
```

When norm level = 2, it is L2 norm:

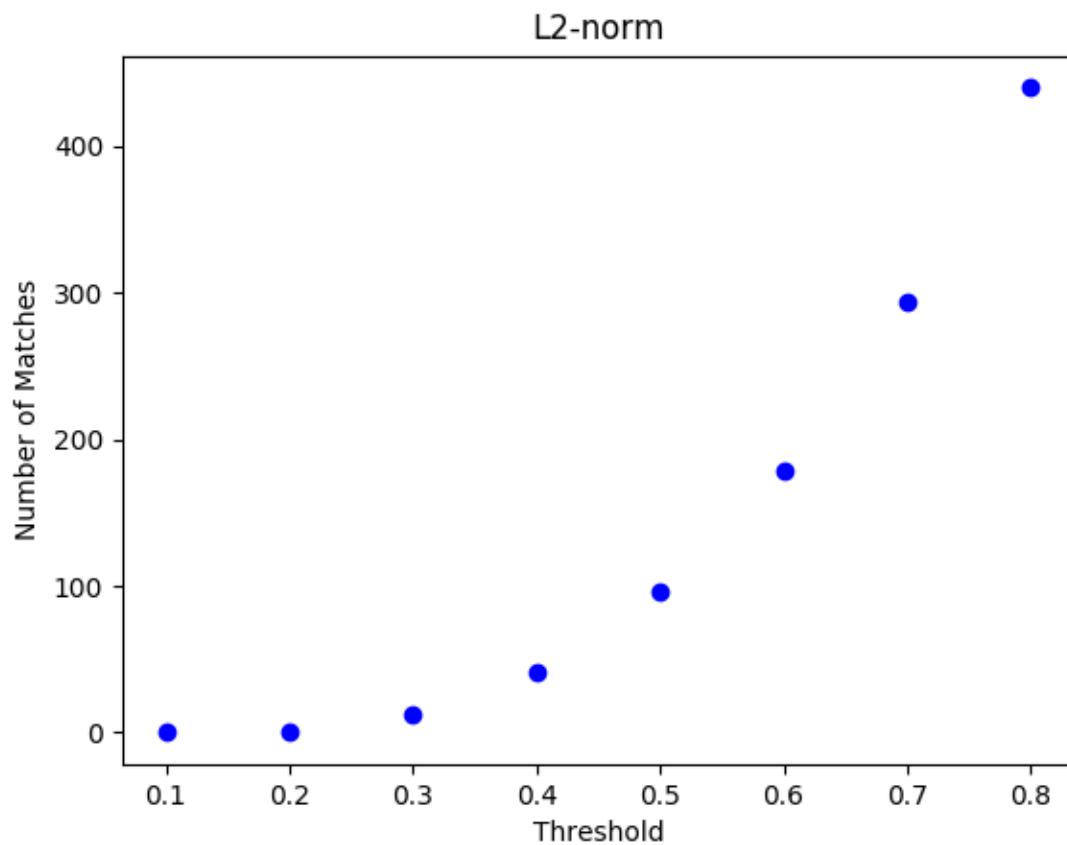
```
458     def euclidean_distance(vector1, vector2, norm_level):
459         vector1 = vector1.reshape(-1)
460         vector2 = vector2.reshape(-1)
461         distance = np.linalg.norm(vector1 - vector2, ord=norm_level)
462         return distance
463
464
465     def show_matching(image1, keypoint1, image2, keypoint2):
466         ax1 = plt.subplot(1, 2, 1)
467         ax2 = plt.subplot(1, 2, 2)
468         img1 = cv2.drawKeypoints(image1, keypoint1, None)
469         img2 = cv2.drawKeypoints(image2, keypoint2, None)
470         ax1.imshow(img1)
471         ax2.imshow(img2)
472
473         colors = ["white", "LightPink", "DarkTurquoise", "PeachPuff", "MediumSlateBlue", "Gold"]
474         color_i = 0
475         for kp1, kp2 in zip(keypoint1, keypoint2):
476             coord1 = kp1.pt
477             coord2 = kp2.pt
478             c = colors[color_i % 6]
479             con = ConnectionPatch(xyA=coord2, xyB=coord1, coordsA="data", coordsB="data",
480                                   axesA=ax2, axesB=ax1, arrowstyle="-", color=c)
481             ax2.add_patch(con)
482             color_i += 1
483         plt.show()
484
485
486     def generate_top_ten(image1, image2, top_ten):
487         img1 = cv2.imread(image1)
488         img2 = cv2.imread(image2)
489
490         gray1 = cv2.cvtColor(img1, cv2.COLOR_BGR2GRAY)
491         gray2 = cv2.cvtColor(img2, cv2.COLOR_BGR2GRAY)
492
493         keypoint1 = []
494         keypoint2 = []
495
496         for item in top_ten:
497             keypoint1.append(item[0][0])
498             keypoint2.append(item[0][1])
499
500         show_matching(gray1, keypoint1, gray2, keypoint2)
```

## Results when threshold increasing:

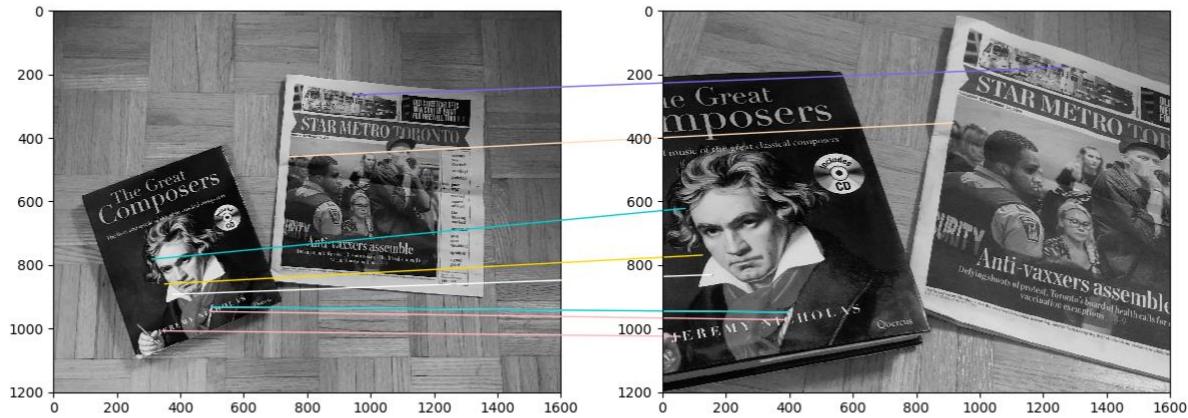


Number of Matches (Please know, since the code execution on my laptop is very long for Q4, to make sure I can finish this assignment on time, I set the limitation as 2000 during sift initialization so that the running can be faster. The results of number of matches and the matching will be slightly less powerful than normal setting. However, the tendency is the same.):

```
360     def sift_extract(image):
361         img = cv2.imread(image)
362         gray_img = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
363
364         # Initialized sift
365         sift = cv2.xfeatures2d.SIFT_create(2000)
366         flag = cv2.DRAW_MATCHES_FLAGS_DRAW_RICH_KEYPOINTS
367
368         # Get keypoints and features
369         keypoints, features = sift.detectAndCompute(gray_img, None)
370         result = cv2.drawKeypoints(gray_img, keypoints, flags=flag, outImage=None)
371
372         cv2.imwrite('sift_keypoints.jpg', result)
373         return keypoints, features
```

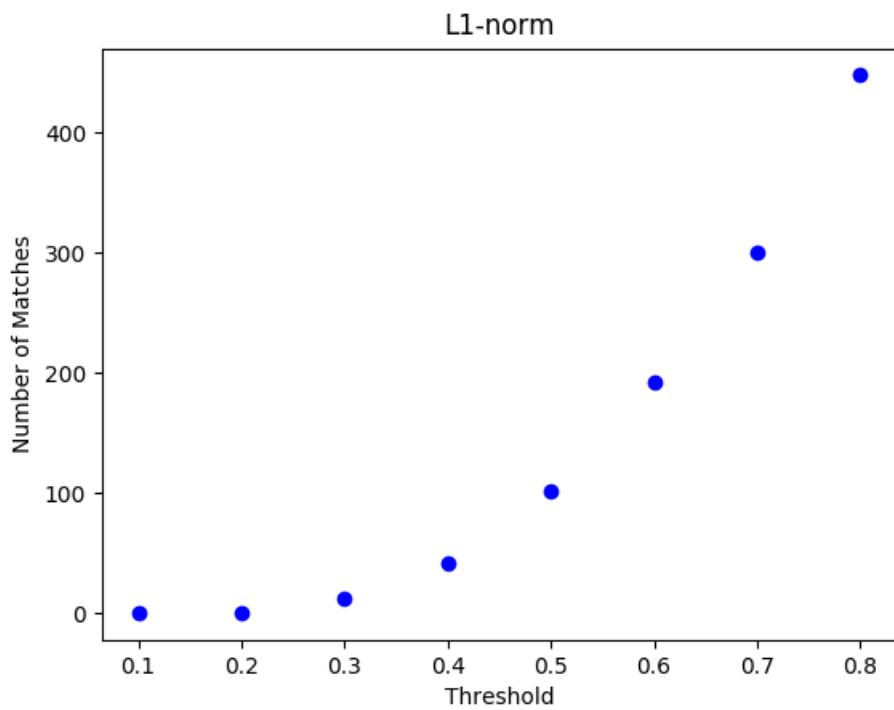


## L2: Top 10 matches at best threshold 0.8

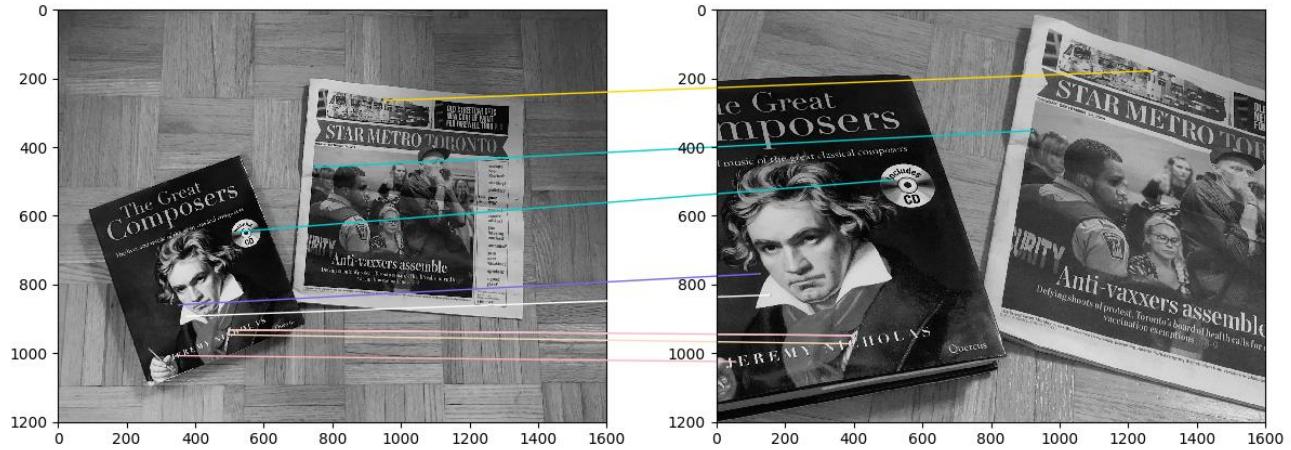


Q4.(c).

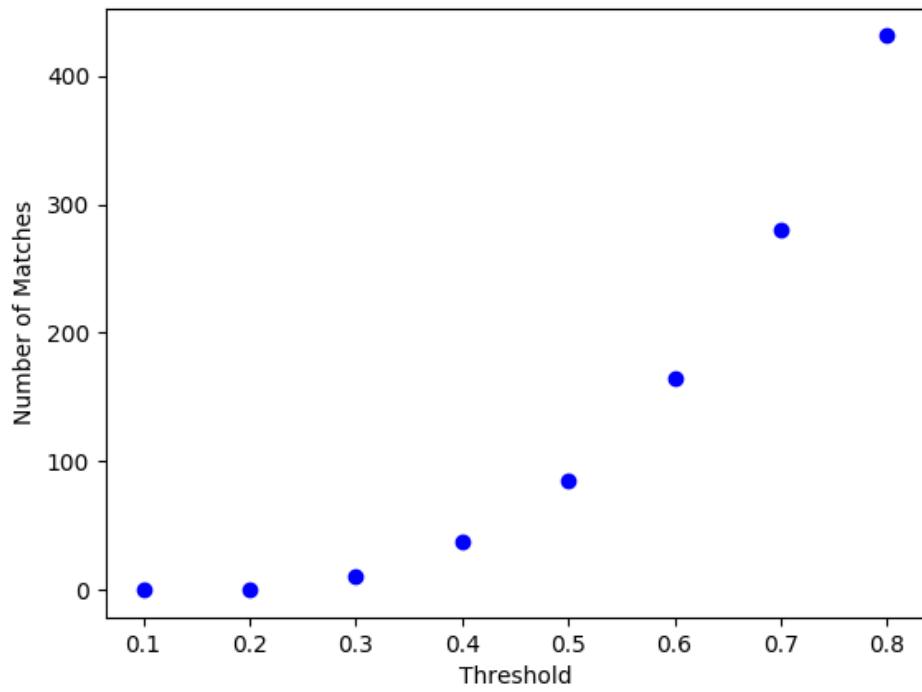
Number of Matches (Please know, since the code execution on my laptop is very long for Q4, to make sure I can finish this assignment on time, I set the limitation as 2000 during sift initialization so that the running can be faster. The results of number of matches and the matching will be slightly less powerful than normal setting. However, the tendency is the same.):



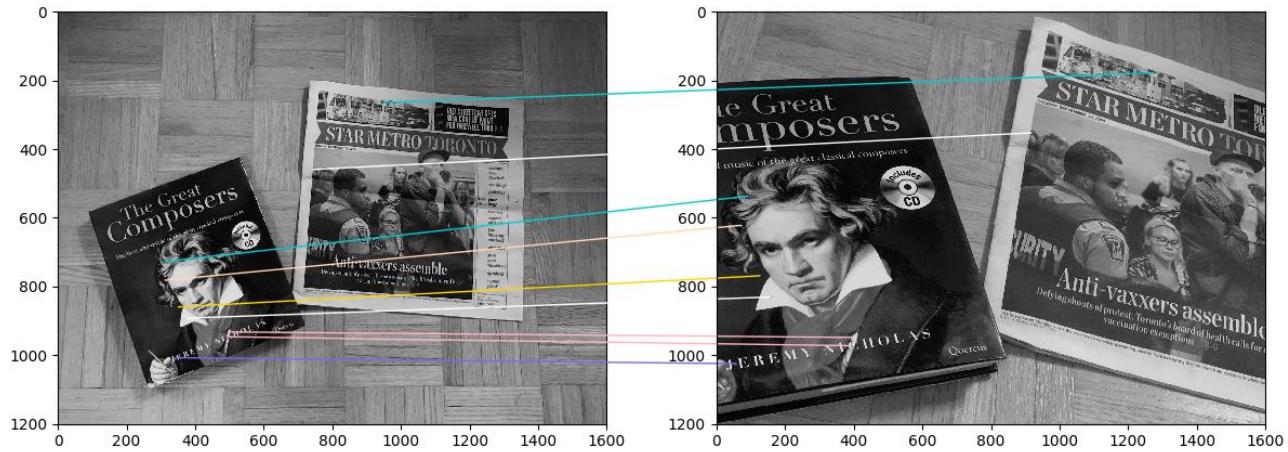
L1: Top 10 matches at best threshold 0.8:



L3-norm



L3: Top 10 matches at best threshold 0.8:

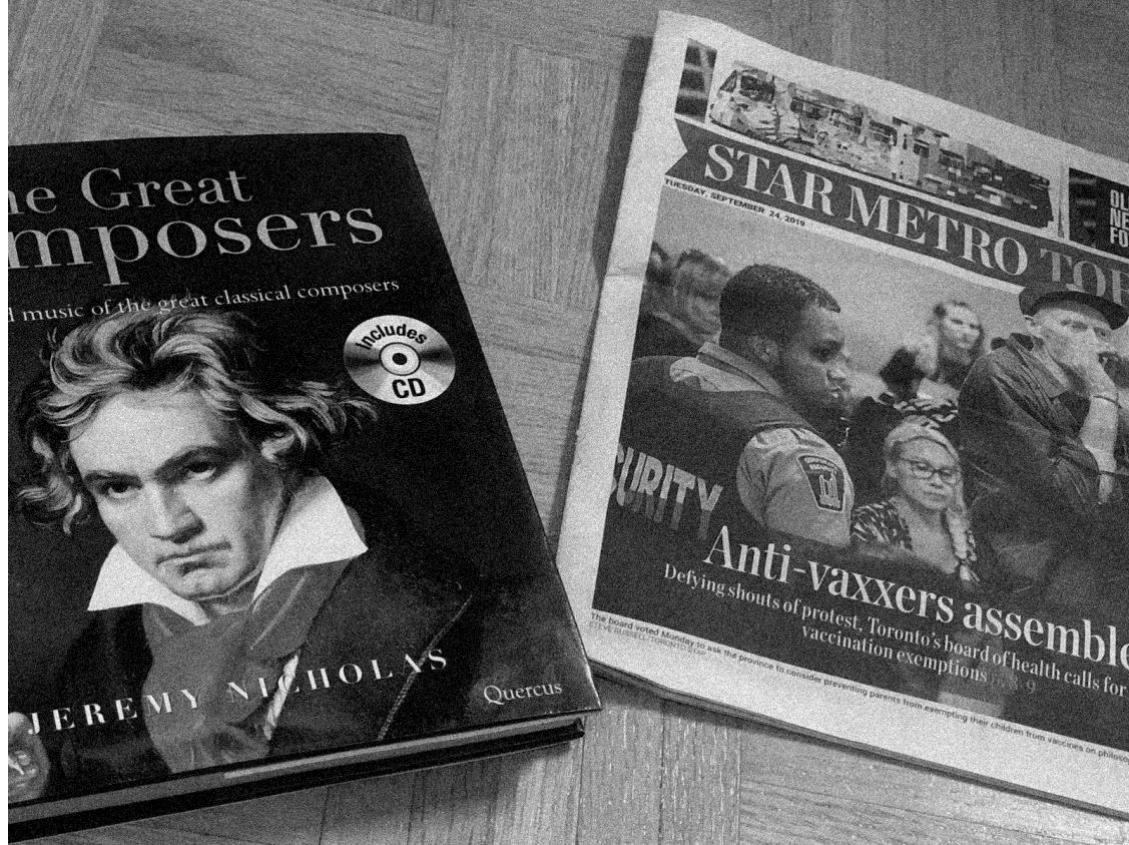
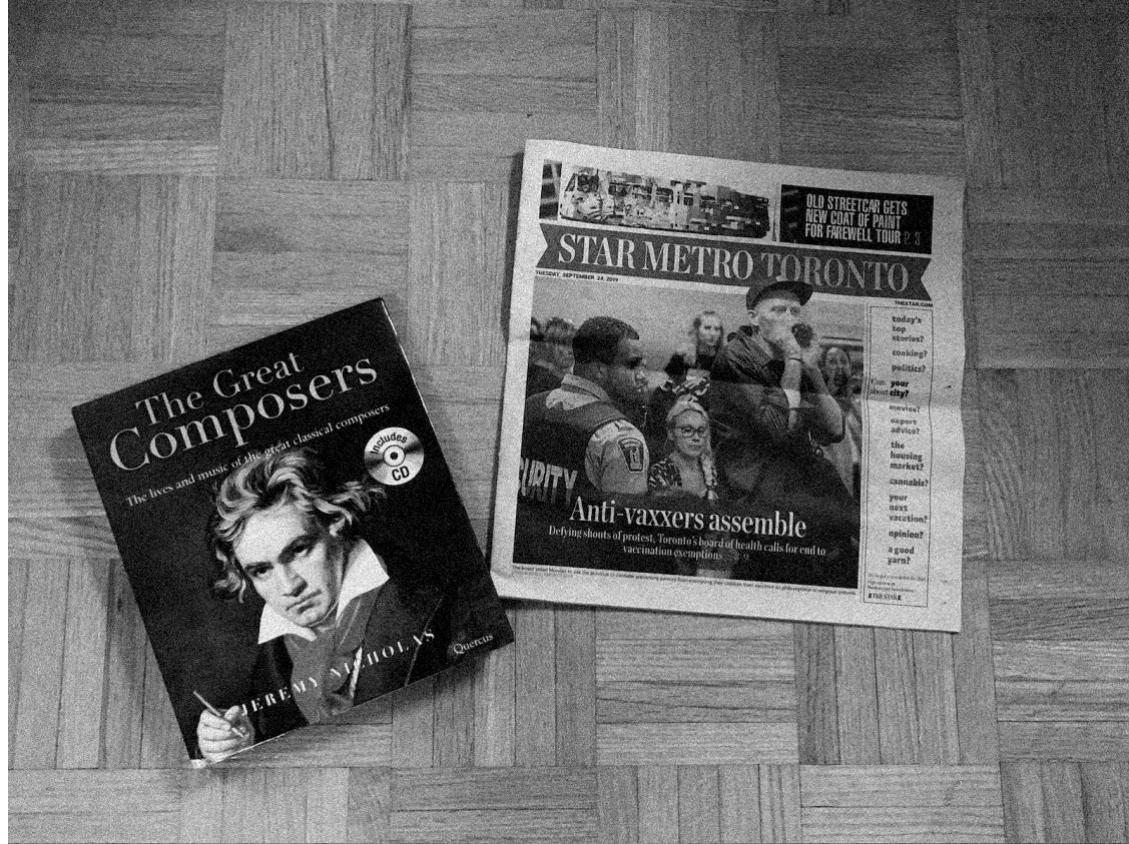


Based on the observation, there is no big difference among L1, L2, L3. We may need more pictures to actually find the differences under varied situations, however, I think the differences won't be large.

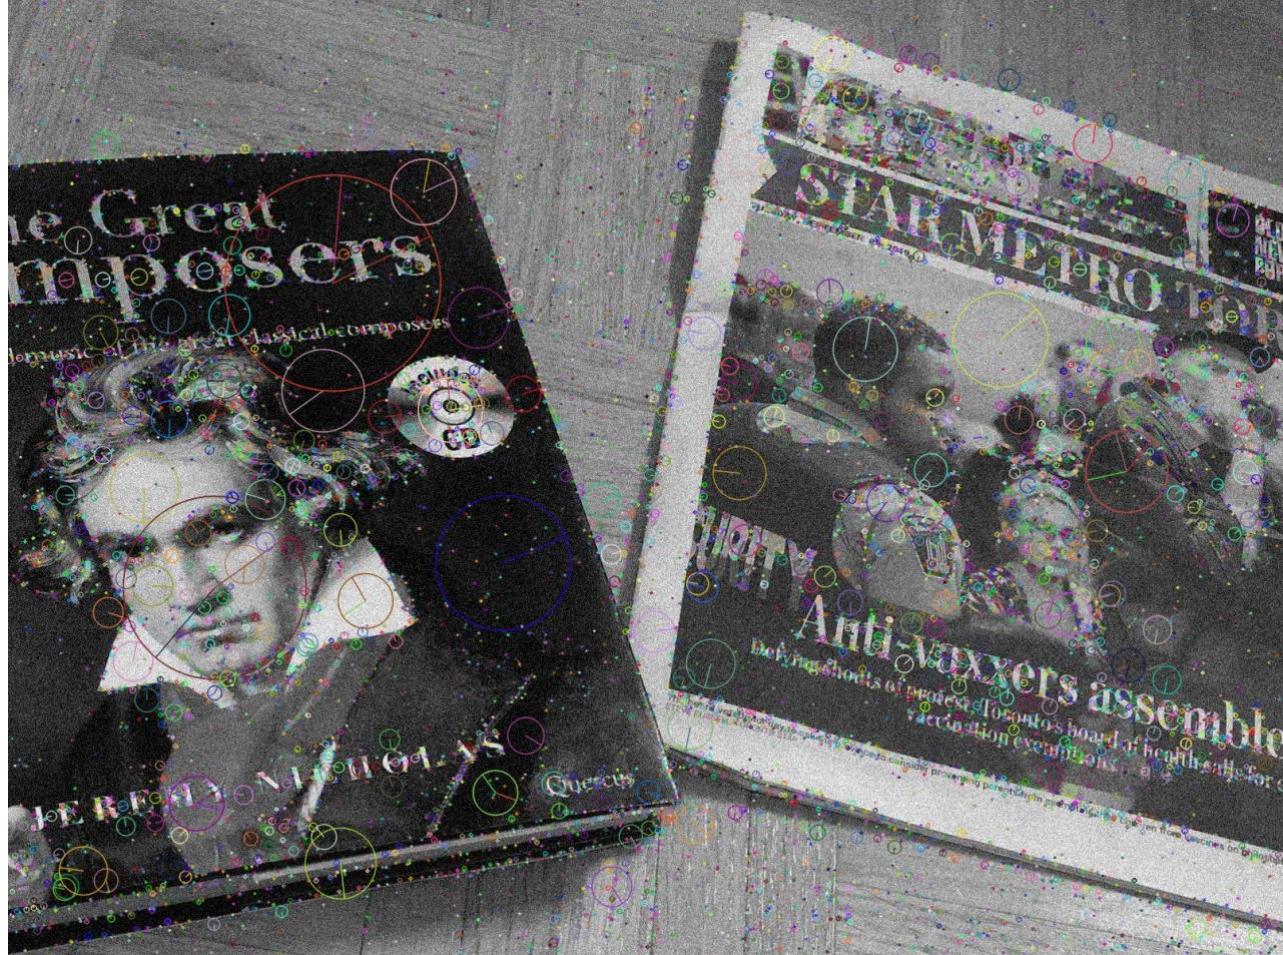
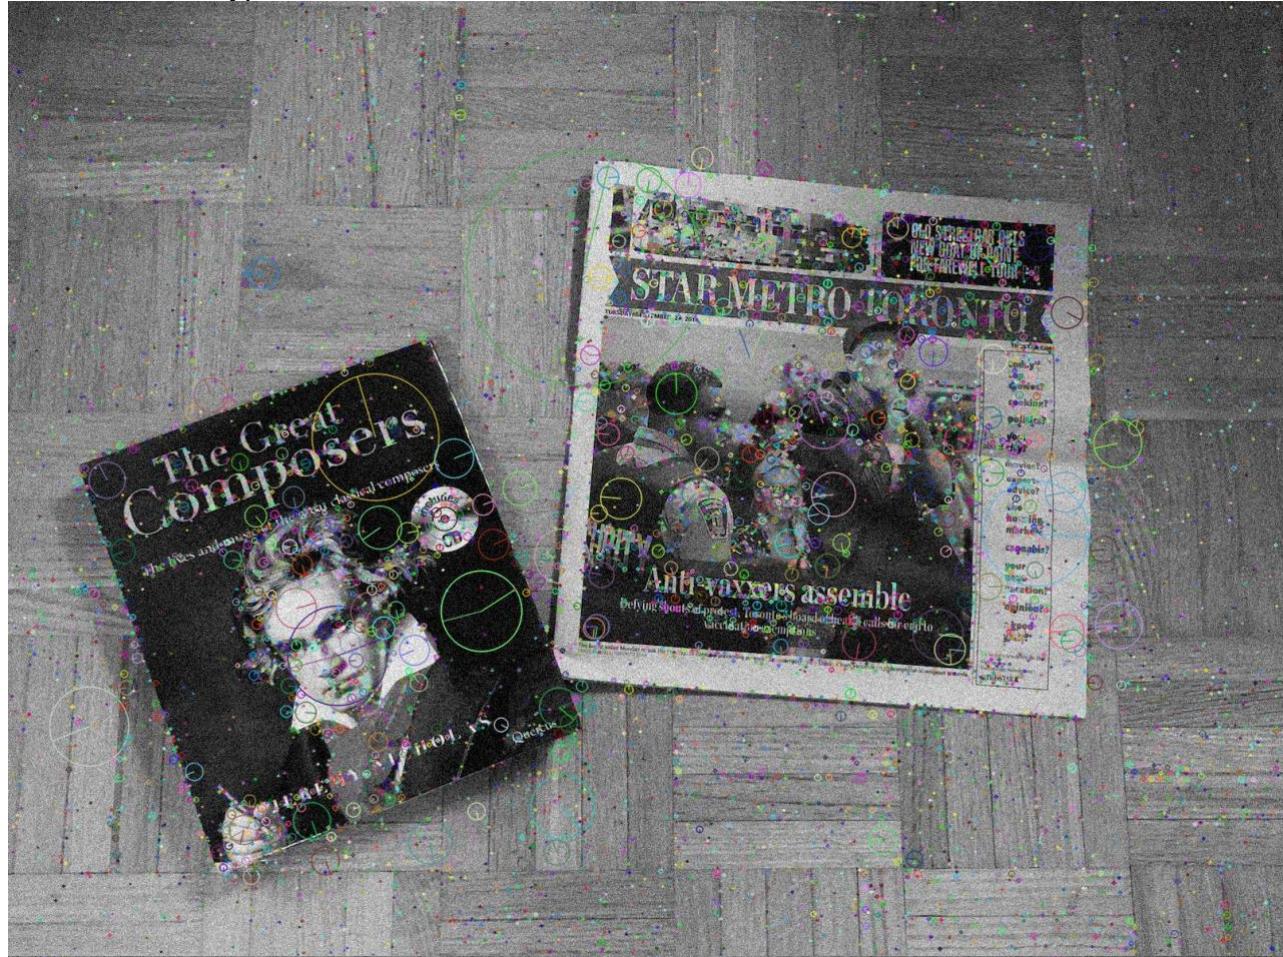
Q4.(d).

```
495 def AddGaussianNoise(image):
496
497     img = cv2.imread(image, cv2.IMREAD_GRAYSCALE)
498     mean = 0
499     sigma = 0.08
500     temp = img / 255
501
502     gaussian = np.random.normal(mean, sigma, size=temp.shape) # np.zeros((224, 224), np.float32)
503
504     noisy_image = temp + gaussian
505
506     return noisy_image * 255
507
```

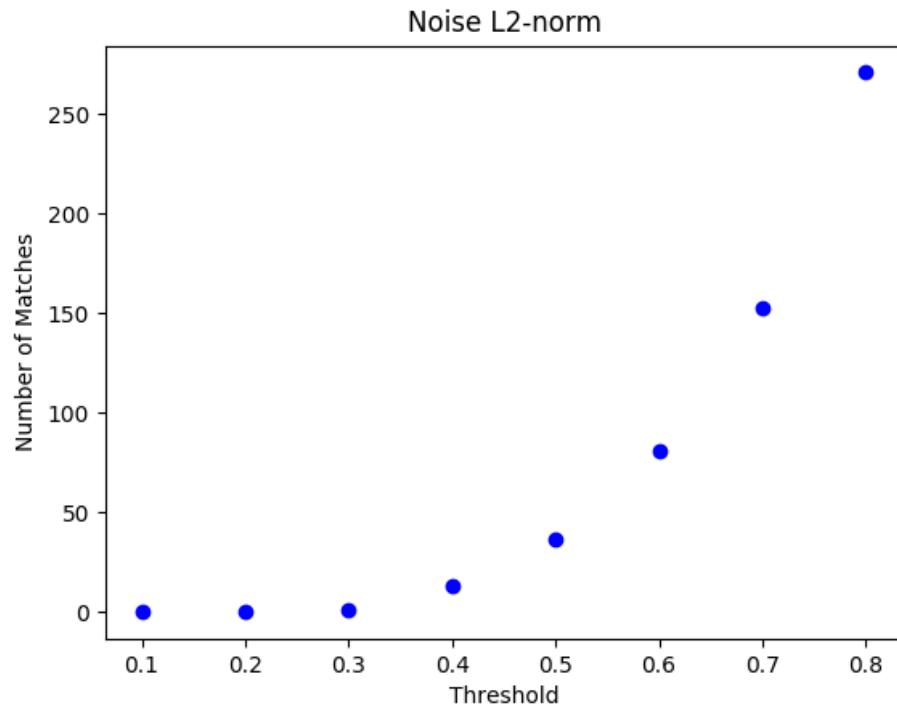
Added noises:



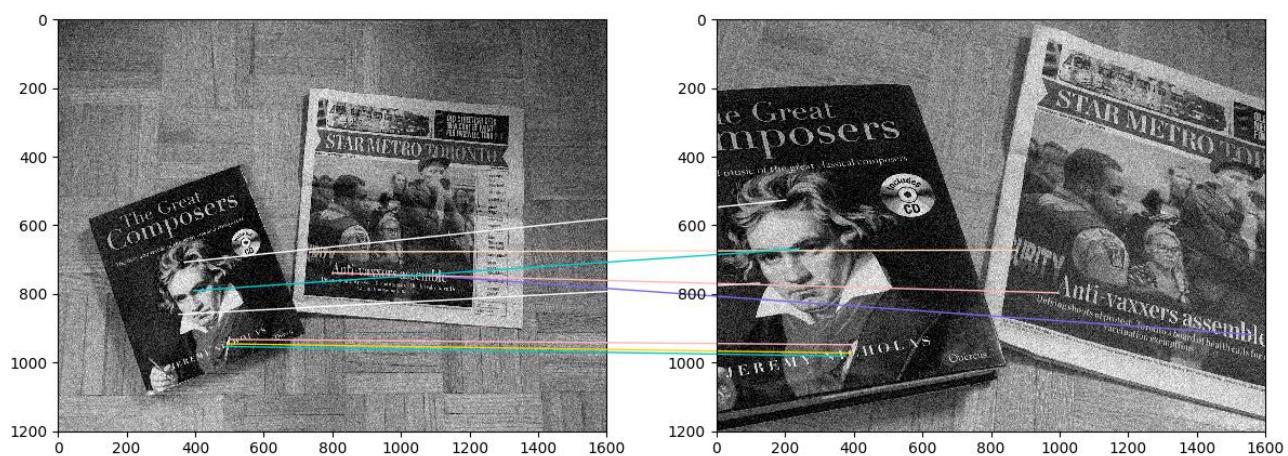
Extract SIFT keypoints:



Number of Matches (Please know, since the code execution on my laptop is very long for Q4, to make sure I can finish this assignment on time, I set the limitation as 2000 during sift initialization so that the running can be faster. The results of number of matches and the matching will be slightly less powerful than normal setting. However, the tendency is the same.):



Top 10 matches at best threshold 0.8:



As you see, the noise decreased the number of matches we found, which is because during the keypoints detection, SIFT mistakenly treat many noises as keypoints, since they are similar to corners. Those wrong keypoints affect the judgement of matching as well. The less number of matches we have, the less matchings we can do.

Q4.(e).

```
509     def sift_matching_color(image1, image2, threshold):
510         img1 = cv2.imread(image1)
511         img2 = cv2.imread(image2)
512
513         # Only use h of hsv (h include all colour information)
514         img1_b = img1[0:img1.shape[0], 0:img1.shape[1], 0]
515         img2_b = img2[0:img2.shape[0], 0:img2.shape[1], 0]
516
517         img1_g = img1[0:img1.shape[0], 0:img1.shape[1], 1]
518         img2_g = img2[0:img2.shape[0], 0:img2.shape[1], 1]
519
520         img1_r = img1[0:img1.shape[0], 0:img1.shape[1], 2]
521         img2_r = img2[0:img2.shape[0], 0:img2.shape[1], 2]
522
523         kp_1, des_1 = sift_extract(image1)
524         kp_2, des_2 = sift_extract(image2)
525
526
527         # Initialize sift
528         sift = cv2.xfeatures2d.SIFT_create(2000)
529
530         kp_1, des_1b = sift.compute(img1_b, kp_1)
531         kp_2, des_2b = sift.compute(img2_b, kp_2)
532         euclidean_dis_b = euclidean_distance2(des_1b, des_2b)
533
534         kp_1, des_1g = sift.compute(img1_g, kp_1)
535         kp_2, des_2g = sift.compute(img2_g, kp_2)
536         euclidean_dis_g = euclidean_distance2(des_1g, des_2g)
537
538         kp_1, des_1r = sift.compute(img1_r, kp_1)
539         kp_2, des_2r = sift.compute(img2_r, kp_2)
540         euclidean_dis_r = euclidean_distance2(des_1r, des_2r)
541
542
543         # find minimum distance index in keypoint 2
544         min_keypoints_index_b = np.argmin(euclidean_dis_b, axis=1)
545         min_keypoints_index_g = np.argmin(euclidean_dis_g, axis=1)
546         min_keypoints_index_r = np.argmin(euclidean_dis_r, axis=1)
547
```

```
548
549     num_matching_keypoint_b = len(min keypoints_index_b)
550     num_matching_keypoint_g = len(min keypoints_index_g)
551     num_matching_keypoint_r = len(min keypoints_index_r)
552
553     # find smallest and 2nd smallest distance
554     smallest_b = np.sort(euclidean_dis_b, axis=1)[:, 0]
555     second_smallest_b = np.sort(euclidean_dis_b, axis=1)[:, 1]
556
557     smallest_g = np.sort(euclidean_dis_g, axis=1)[:, 0]
558     second_smallest_g = np.sort(euclidean_dis_g, axis=1)[:, 1]
559
560     smallest_r = np.sort(euclidean_dis_r, axis=1)[:, 0]
561     second_smallest_r = np.sort(euclidean_dis_r, axis=1)[:, 1]
562
563     # compute ratio
564     ratio_b = np.true_divide(smallest_b, second_smallest_b)
565     ratio_g = np.true_divide(smallest_g, second_smallest_g)
566     ratio_r = np.true_divide(smallest_r, second_smallest_r)
567
568     keypoint1 = []
569     keypoint2 = []
570     matched_ratio = []
571     keypoints_pairs = {}
572     count = 0
573
574     for i in range(num_matching_keypoint_b):
575         print(ratio_b[i])
576         if ratio_b[i] < threshold:
577             count = count + 1
578             kp1 = kp_1[i]
579             kp2 = kp_2[min keypoints_index_b[i]]
580             keypoint1.append(kp1)
581             keypoint2.append(kp2)
582             matched_ratio.append(ratio_b[i])
583             keypoints_pairs[(kp_1[i], kp2)] = ratio_b[i]
```

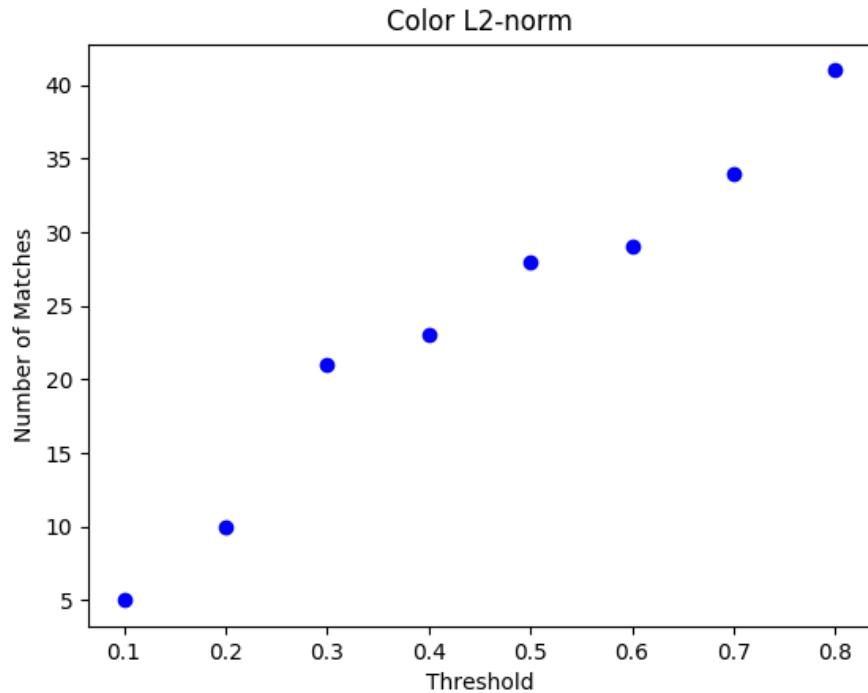
```

585     for i in range(num_matching_keypoint_g):
586         if ratio_g[i] < threshold:
587             count = count + 1
588             kp1 = kp_1[i]
589             kp2 = kp_2[min_keypoints_index_g[i]]
590             keypoint1.append(kp1)
591             keypoint2.append(kp2)
592             matched_ratio.append(ratio_g[i])
593             keypoints_pairs[(kp_1[i], kp2)] = ratio_g[i]
594
595     for i in range(num_matching_keypoint_r):
596         if ratio_r[i] < threshold:
597             count = count + 1
598             kp1 = kp_1[i]
599             kp2 = kp_2[min_keypoints_index_r[i]]
600             keypoint1.append(kp1)
601             keypoint2.append(kp2)
602             matched_ratio.append(ratio_r[i])
603             keypoints_pairs[(kp_1[i], kp2)] = ratio_r[i]
604
605     sorted_keypointes = sorted(keypoints_pairs.items(), key=lambda kv: kv[1])
606
607     print("number of matching = ", count)
608
609     top_ten = sorted_keypointes[0:9]
610     print(top_ten)
611
612     flag = cv2.DRAW_MATCHES_FLAGS_DRAW_RICH_KEYPOINTS
613
614     result_img1 = cv2.drawKeypoints(img1, keypoint1, outImage=None, flags=flag)
615     result_img2 = cv2.drawKeypoints(img2, keypoint2, outImage=None, flags=flag)
616     cv2.imwrite("sample1_keypoints_final_color.jpg", result_img1)
617     cv2.imwrite("sample2_keypoints_final_color.jpg", result_img2)
618     show_matching(img1, keypoint1, img2, keypoint2)
619
620     return top_ten
621
622
623
624
625     def euclidean_distance2(vector1, vector2):
626         vector1_square = np.sum(np.square(vector1), axis=1).reshape(-1, 1)
627         vector2_square = np.sum(np.square(vector2), axis=1).reshape(-1, 1)
628         vector1_vector2 = np.dot(vector1, np.transpose(vector2))
629
630         distance = np.sqrt(vector1_square - 2 * vector1_vector2 +
631                            np.transpose(vector2_square))
632
633         return distance

```

The algorithm is similar to Q4b, but with extra R, G, and B channels. Keypoints are extracted from grayscale image, then compute three descriptors separately from R, G, B channels. Then perform sift match on these three descriptors to calculate new matched keypoints.

Number of Matches (Please know, since the code execution on my laptop is very long for Q4, to make sure I can finish this assignment on time, I set the limitation as 2000 during sift initialization so that the running can be faster. The results of number of matches and the matching will be slightly less powerful than normal setting. However, the tendency is the same.):



Top 10 matches at best threshold 0.8:

